


training (NLL)
make_train_sequences



evaluation (forward)
make_eval_sequences



input ends at $i+191$, target begins at $i+192$: disjoint and consecutive



Both rows are the index arrays the two builders return for window i , drawn to one shared scale. At evaluation the whole input is passed as past_target with past_observed_target all ones, and no target value is anywhere in it. Training's two halves are concatenated into a single $96+h$ tensor, but `_val_loss` splits it at `context_length` and the module substitutes its learned `mask_encoding` at every target position, so the target is scored by the likelihood and never conditioned on.