
Drift–Utility Dissociation in Time-Series Foundation Model Fine-Tuning: Methodology and a Moirai Case Study

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We contribute a **reusable diagnostic methodology** for assessing representation
2 drift during fine-tuning of time-series foundation models, comprising a value-
3 gate screening criterion, controlled-comparison design, and task-orthogonal probe
4 protocol. On Moirai/ETTh2 (gate-pass 28%), task-orthogonal probe asymme-
5 try ($\Delta R^2 > 0$ on the trained head, < 0 on positive-floor probes; 10/10 seeds) es-
6 tablishes that fine-tuning restructures toward the objective rather than improving
7 generally. Restructuring is generally beneficial for gate-passing cells; harm oc-
8 curs only in the narrow low- n /small-capacity regime (1 of 7 cells tested). Cross-
9 backbone validation on Chronos-T5-Small/M4-Monthly (gate-pass 84.5%) con-
10 firms the value-gate as a regime indicator: strong gate-pass predicts encoder sta-
11 bility ($\text{CKA} \geq 0.95$), no probe signal, and no forgetting. On ILI (gate-pass 57%),
12 full probes show restructuring with beneficial outcome ($\text{CKA} = 0.478$, $\Delta R^2 > 0$ in
13 10/10 seeds, forgetting -80.7%), extending the probe-asymmetry regularity to a
14 third gate-passing cell and revising the predicted stability threshold upward. **Prac-**
15 **titioner takeaway:** apply the value-gate before investing in drift diagnosis; use
16 LoRA with reduced LR on larger variants.

17 1 Introduction

18 Moirai Woo et al. [2024] is a foundation time-series model trained on billions of time points that
19 achieves strong zero-shot performance on forecasting benchmarks. A natural next step is to fine-
20 tune it on domain-specific tasks—but when should a practitioner worry about overwriting the rep-
21 resentations that made zero-shot work? *Counterfactual:* Chronos-T5-Small on ETTh2 gives ZS
22 MSE=0.304 vs. Linear 0.213 (ZS *underperforms* by 43%). A practitioner who skipped screen-
23 ing would fine-tune Chronos, observe CKA dropping, and apply LoRA/EWC to preserve features
24 that had *no advantage over Linear*—wasted effort. Our value-gate prevents this: drift diagnosis is
25 meaningful only when pre-trained features demonstrably outperform a trivial baseline.

26 **Representation drift during fine-tuning.** We contribute a systematic *diagnostic methodology* for
27 characterising representation drift—progressive overwriting of pre-trained features—during fine-
28 tuning of time-series foundation models, and demonstrate it on Moirai (Small 13.8M, Base 91M,
29 Large 311M) as the single backbone family where pre-trained features pass our value-gate on ETT
30 benchmarks. Matched-protocol attempts on three non-Moirai backbones (Chronos-T5-Small Ansari
31 et al. [2024], TimesFM-2.5-200M Das et al. [2024], MOMENT-1-base Goswami et al. [2024]) gate-
32 fail on all nine ETT cells tested—the expected behaviour of a well-calibrated screening criterion,
33 since these corpora lack ETT coverage. Chronos gate-passes M4-Monthly (in its corpus) at 84.5%,
34 confirming the gate identifies applicable cells across backbones (§5.1). Catastrophic forgetting dur-
35 ing fine-tuning is well-studied in continual learning Kirkpatrick et al. [2017], Li and Hoiem [2018]

and NLP Luo et al. [2023]; our contribution is reusable diagnostic methodology, demonstrated on a well-controlled Moirai case study. We use “catastrophic forgetting” to mean degraded performance on distributions the model originally handled well—consistent with continual learning usage where the “previous task” is the pre-training objective.

Distinguishing drift from overfitting. Representation drift degrades zero-shot performance while training loss decreases (not overfitting); the frozen-encoder control (CKA=1.000, same data/epochs) isolates encoder weight updates as the cause.

On ETTh2 (gate-pass: Moirai outperforms Linear by 28–45%, $R^2_{\text{task}}(\text{PT})=+0.279$), fine-tuning degrades MSE by 10–16% at $n=500$ and resolves by $n=2,000$. ETTh1 and ETTm2 gate-fail at the standard evaluation protocol and are included as cross-domain comparisons (Table 3). On IoT (sub-random ZS AUC), CKA drops to 0.31–0.45 but task performance improves from a near-null baseline: the drift is irrelevant because there were no valuable features to forget.

Why this matters. Time-series drift directly harms distributional predictions for anomaly scoring and probabilistic forecasting. Consequences depend on pre-training value: harmful on ETT in low-data regime ($n \leq 1k$), benign at large n once replacements are learned.

Scope. The empirical case study spans three gate-passing cells: Moirai-Small/Base on ETTh2, Moirai-Small on ILI, and Chronos-T5-Small on M4-Monthly. The central empirical finding is **non-monotone capacity dependence**: Base (91M) improves protocol-robustly at all n ; Small (13.8M) degrades at low n ; Large (311M) degrades at default LR but is rescued by $10\times$ LR reduction (forg. -8.5% , CKA=0.989; Appendix F). A step-size falsification test (Base at $10\times$ LR, 10 seeds) restructures *more* aggressively than Small (CKA=0.14 vs. 0.94) yet still improves (-58.2% ; Appendix G)—capacity provides an independent buffer against destructive drift. Both Small and Base share the probe-asymmetry mechanism (10/10, 9/10 seeds), confirming that drift harm depends on capacity relative to the adaptation target, not on the restructuring itself. We do not claim a universal phenomenon; we claim reusable methodology demonstrated across three (backbone, dataset) configurations. The methodology transfers to any backbone that gate-passes; the specific numeric findings require broader replication.

Experimental approach. We deploy the diagnostic methodology across three Moirai sizes (Small/Base/Large) on ETT forecasting Zhou et al. [2021] (ETTh2 gate-pass; ETTh1/ETTM2 cross-domain) and IoT anomaly detection (negative control; §5). Forgetting sign at $n=10k$ is protocol-dependent (-5.3% ES vs. $+7.5\%$ final-epoch); we lead on the more robust $\Delta R^2 > 0$ majority.

Contributions.

1. **Diagnostic methodology (backbone-agnostic).** Value-gate screening (ZS $> 20\%$ vs. Linear), controlled-comparison design (frozen-encoder, random-init, layer-wise unfreeze), and task-orthogonal probe protocol. Supporting infrastructure (uni2ts patch, CUDA-deterministic protocol, LoRA-Large $10\times$ LR recipe) released with code.
2. **Empirical validation across three gate-passing cells.** Probe asymmetry on Moirai/ETTh2 (10/10 seeds), ETTm2 (10/10), and ILI (10/10; gate-pass 57%, CKA=0.478; §5.1); encoder-stability confirmation on Chronos/M4 (gate-pass 84.5%, CKA ≥ 0.95 ; §5.1).

2 Background and Problem Setup

Catastrophic forgetting in fine-tuning. Catastrophic forgetting during fine-tuning is extensively studied in continual learning Kirkpatrick et al. [2017], Li and Hoiem [2018] and NLP Luo et al. [2023]. Kumar et al. [2022] show full fine-tuning can distort pre-trained features relative to linear probing; Andreassen et al. [2023] characterize OOD robustness decay throughout fine-tuning. Our contribution is a systematic *characterization* for time-series foundation models: CKA, linear probing, weight drift, and frozen-encoder controls across domains with contrasting pre-training value.

Time-series foundation model adaptation. Recent work on adapting pre-trained time-series models includes probabilistic fine-tuning (Lag-Llama Rasul et al. [2024]), generative pre-training with task-specific heads (Timer Liu et al. [2024]), LLM reprogramming (Time-LLM Jin et al. [2024]), and cross-task transfer (One Fits All Zhou et al. [2023]). These works develop *adaptation recipes*;

our contribution is orthogonal: a *diagnostic methodology* for characterising what happens to pre-trained representations during any such adaptation, applicable when the practitioner needs to decide whether drift is harmful or benign.

Setting. Given a pre-trained time-series foundation model with encoder $f_\theta : \mathbb{R}^{T \times F} \rightarrow \mathbb{R}^d$ and downstream dataset $\mathcal{D}$ of size n , fine-tuning produces updated parameters θ' by minimising a task loss (NLL for forecasting). We measure encoder changes using three diagnostics (§3): (1) CKA between f_θ and $f_{\theta'}$ (geometric similarity), (2) linear probe R^2 onto the downstream target ($\Delta R^2 = R^2(\text{FT}) - R^2(\text{ZS})$), and (3) weight drift $\|\theta' - \theta\|_2$.

Pre-training value gate. Studying whether fine-tuning destroys valuable features requires a domain where such features are demonstrably present. We set the gate at zero-shot MSE 20% below a Linear baseline; ETTh2 at $h=96$ passes with Moirai-Small ($R^2_{\text{task}}(\text{PT}) = +0.279$). ETTh1 and ETTm2 gate-fail at the standard evaluation protocol ($R^2_{\text{task}}(\text{PT}) = -0.41$ and -0.56 respectively; Table 3); they are included as cross-domain comparisons to characterise drift behaviour outside the gate-pass regime. Domains that gate-fail serve as negative controls. The 20% threshold is calibrated to ETT short-horizon benchmarks where linear baselines are surprisingly competitive Zeng et al. [2023]: a model must substantially outperform a lookback-96 linear regression before drift can be *harmful* (rather than irrelevant). The threshold is a practical screen, not a universal constant; we report $R^2_{\text{task}}(\text{PT})$ values for each cell (Table 3) so readers can apply alternative thresholds.

Task-level consequence: forgetting. We define forgetting as signed relative change in zero-shot-canonical task metric on held-out test: $\text{forg.\%} = 100 \cdot (M_{\text{FT}} - M_{\text{ZS}}) / |M_{\text{ZS}}|$ (MSE for forecasting, so positive = worse). $\text{Forg.} > 0$ indicates catastrophic forgetting (degraded performance vs. pre-training baseline).

3 Fine-Tuning Protocol

Moirai for forecasting and anomaly scoring. Moirai Woo et al. [2024] predicts a probabilistic distribution over a target horizon given context. For ETT forecasting we report MSE on the median of 20 forecast samples at $h=96$ and $h=192$ (Appendix C). For IoT the anomaly score is the fraction of timesteps outside the $1-\alpha$ CI ($\alpha=0.05$) over a 32-timestep horizon.

Fine-tuning objective. Condition B uses NLL only. Condition C adds supervised contrastive loss Khosla et al. [2020]: $\mathcal{L}(\theta) = \mathcal{L}_{\text{NLL}}(\theta) + \lambda \mathcal{L}_{\text{SupCon}}(\mathbf{z}, \mathbf{y}; \tau)$, with $\lambda = 0.5$ and $\tau = 0.07$. For ETT, temporal clusters (8 bins by time-series position) serve as SupCon label.

Training details. AdamW ($\eta = 10^{-4}$, weight decay 10^{-2}), batch size 32, early stopping with patience 3, gradient clipping at 1.0. ETT fine-tuning runs 20 epochs at $n=500$ (10 seeds) and is swept across $n \in \{500, 1k, 2k, 5k, 10k\}$. All Moirai sizes (Small 13.8M, Base 91M, Large 311M) use the same protocol.

Deterministic training protocol. All $n=10k$ per-seed results use CUDA-deterministic training: `CUBLAS_WORKSPACE_CONFIG=:4096:8`, `torch.backends.cudnn.deterministic=True`, seeded DataLoader workers, and fixed `PYTHONHASHSEED`. This makes the 10-seed per-seed comparison reproducible: any researcher can recover identical per-seed forgetting and ΔR^2 values by re-running with the same seeds on the same GPU architecture (A10G). Full seed list, per-seed values, and hardware specification are in Appendix I. Cross-hardware reproducibility (e.g., A100 vs. A10G) has not been empirically verified; the deterministic guarantee applies within the same GPU architecture (A10G, as used in this paper).

uni2ts PackedStdScaler patch.¹ Any Moirai fine-tuning reproduction of this paper requires the patch.

Representation diagnostics. We track representation change with three metrics: (1) CKA Kornblith et al. [2019] between pre-trained and fine-tuned encoder representations (geometric similarity); (2) linear probing R^2 (Ridge, MLP, Linear-Forecaster heads on frozen representations, predicting forecasting target; we report $\Delta R^2 = R^2(\text{FT}) - R^2(\text{ZS})$; absolute $R^2(\text{FT})$ remains negative, so ΔR^2

¹We patched an in-place tensor mutation in uni2ts’s `PackedStdScaler` that silently detaches autograd; a differentiable `torch.where` replaces the in-place write. All results use the patched scaler (released with our code). See Appendix B for details.

quantifies relative improvement over pre-trained linear decoding); and (3) weight drift (ℓ_2 distance from pre-trained parameters).

Table 1: Notation summary for probe-related quantities.

Symbol	Definition
$R_{\text{task}}^2(\text{PT})$	$1 - \text{MSE}_{\text{ZS}}/\text{MSE}_{\text{Linear}}$; positive $\Rightarrow$ gate-pass
$R^2(\text{PT}), R^2(\text{FT})$	Ridge probe score on pre-trained / fine-tuned encoder
ΔR^2	$R^2(\text{FT}) - R^2(\text{PT})$; relative linear-decodability change
ΔR_{task}^2	Task-level improvement measured via forecasting head

4 Representation Drift Diagnosis: ETT Forecasting

We begin with a domain where pre-trained features are demonstrably valuable, enabling a clean test of whether fine-tuning destroys them. ETT time-series forecasting benchmarks Zhou et al. [2021] provide this setting: zero-shot Moirai outperforms Linear baselines by 28–45% (Section 4.1). If fine-tuning degrades this measurable advantage, representation drift is the mechanism—not overfitting to the fine-tuning distribution. The frozen-encoder control (same data, same epochs, CKA=1.000) confirms this: overfitting would not be eliminated by freezing.

4.1 Experimental Setup and Value Gate

We evaluate on ETTh2 (7 features, 17k timesteps, 12/4/4 month split). Moirai-Small (13.8M) zero-shot outperforms Linear by 28–45% (ZS MSE: 0.126 at $h=96$, 0.158 at $h=192$; un-normalised gate-pass scale: 0.269 vs. 0.373 at $h=96$, 0.299 vs. 0.544 at $h=192$), passing the value gate (20% threshold). The 20% threshold is a practical screen, not a phase boundary: at 10–20%, drift diagnosis applies but marginal features reduce both harm from drift and benefit from mitigation; we recommend full probes at $\geq 20\%$ and CKA-only at 10–20%. Conditions: **B** (NLL-only), **C** (NLL+SupCon), **D** (frozen encoder), **E** (LoRA $r=8$), **F** (L2-SP), **G** (EWC). Training: $n=500$ samples, 20 epochs, AdamW (lr=1e-4), 10 seeds (B/C/D) or 5 seeds (E/F/G). Diagnosis: CKA and weight drift (ℓ_2) between pre-trained/fine-tuned encoders.

4.2 Representation Drift Degrades High-Value Features

Fine-tuning degrades ETTh2 performance where pre-trained features are valuable (Appendix Table 5). Conditions B (NLL-only) and C (NLL+SupCon) degrade by 10–16% MSE (10/10 seeds positive, $p=0.002$ sign test; $d=1.23$, 95% CI [0.37, 2.01]; bootstrap $p<10^{-4}$) while training loss *decreases*—not overfitting (which raises validation loss), but representation drift. The frozen-encoder control (D, CKA=1.000, same data/epochs) eliminates drift, ruling out the training regime. Degradation scales with horizon: worse at $h=192$ (13–16%) than $h=96$ (10–11%). Contrastive loss (C) does not prevent drift: C (+10.1%) and B (+11.1%) differ by 1.0%, within seed standard deviation—a null result for the contrastive objective.

4.3 Mitigation Spectrum

We evaluate five mitigation strategies (Appendix Table 6): LoRA (E), frozen encoder (D), EWC (G), L2-SP (F), and contrastive loss (C). Methods range from unconstrained fine-tuning (most drift) to full encoder freezing (no drift).

LoRA achieves the best trade-off. LoRA ($r=8$, $\alpha=16$) *improves* over zero-shot at $h=96$ while maintaining high CKA (0.979) and zero weight drift. On Moirai-Large, the same recipe fails at default LR= 10^{-4} ($+25.7\pm 3.2\%$, 3 seeds); rank escalation does *not* rescue. Reducing LR by $10\times$ to 10^{-5} does rescue: $\text{forg.} = -8.5\pm 0.7\%$ (CKA=0.989, 5 seeds), transferring cross-dataset to ETTh1/ETTh2 at $n=500$. **Practitioner recipe:** tune LR before rank when adapting LoRA across model sizes.

EWC and L2-SP reduce but do not eliminate drift. Both improve CKA over unconstrained fine-tuning but MSE remains 18–19% above zero-shot. Frozen encoder (D) eliminates drift (CKA=1.000) but cannot adapt—a limitation when pre-trained features are insufficient.

Table 2: Sample-size sweep on ETTh2 ($h=96$, condition B). Drift grows monotonically (CKA $\downarrow$ as $n\uparrow$); forgetting is *non-monotonic*: severe at $n\leq 1,000$, resolved at $n=2,000$, with elevated variance at $n=5,000$ (crossover regime). At $n=10,000$: CUDA 10-seed deterministic early-stopped, 7/10 negative forgetting; **10/10** $\Delta R^2 > 0$ (per-seed values in Appendix I). ΔR^2 is a *relative* signal: absolute R^2 (FT) remains negative ($R_{ZS}^2 = -6.90$ for all rows). Ridge ΔR^2 at $n=500$ ($k=2$) and $n=1,000$ ($k=1$): point values only.

Samples	MSE	Forg.%	CKA	R^2 (FT)	ΔR^2 (Ridge)	ΔR^2 (MLP)
500	.148 $\pm$.008	+14.0 $\pm$ 6.7	.949 $\pm$.010	-6.79 ($k=2$)	+12 ($k=2$)	-.26 ($k=1$)
1,000	.147 $\pm$.010	+13.4 $\pm$ 9.1	.936 $\pm$.025	-6.77 ($k=1$)	+13 ($k=1$)	+30 ($k=1$)
2,000	.126 $\pm$.008	-2.9 $\pm$ 5.9	.814 $\pm$.042	-6.51 ($k=3$)	+39 $\pm$.35	+73 $\pm$.75
5,000	.134 $\pm$.012	+6.6 $\pm$ 10.0	.546 $\pm$.082	-6.41 ($k=7$)	+50 $\pm$.26	—
10,000 ^{cuda}	.122 $\pm$.008	-5.3 $\pm$ 6.2	.518 $\pm$.188	—	+67 $\pm$.23 [†]	—

^{cuda} CUDA 10-seed deterministic (A10G); all completed, mean best epoch 4.2 $\pm$ 1.4 (Appendix I for per-seed values and protocol comparison). [†]**10/10** positive; 95% Wilson CI [0.72, 1.00]. Forgetting sign: 7/10 negative (95% Wilson CI [0.40, 0.89]; protocol-dependent—see §5). $\Delta R^2 > 0$ is the robust signal.

4.4 Sample-Size Sweep: Drift–Forgetting Dynamics

Table 2 ($n \in \{500, \dots, 10k\}$, ETTh2) reveals that *drift and forgetting are dissociable*—a finding best illustrated by comparing Moirai-Base (protocol-robust improvement of 50–57%, 10/10 seeds; probe asymmetry 9/10 $\Delta R^2 > 0$ despite CKA $\approx$ 0.25; §5) with Small, where the same mechanism produces capacity-dependent outcomes. On Small, drift grows monotonically with n (CKA from 0.949 at 500 to 0.518 at 10k), but forgetting is non-monotonic: severe at low n (+14% at 500), resolved by 2k (−2.9%), predominantly negative at 10k (−5.3 $\pm$ 6.2%, 7/10 negative, 10-seed CUDA deterministic). The $n=5k$ variance (+6.6 $\pm$ 10.0%, 7 seeds) is bimodal: 4 seeds cluster at −0.1%, 3 at +15.5% with epoch-1 overshoot (full trajectory analysis in Appendix I.1). Linear probing shows ΔR^2 increasing with n : +0.12 ($k=2$ probe seeds) at 500 to +0.67 $\pm$ 0.23 at 10k (**10/10** seeds positive). Geometric dissimilarity (CKA $\downarrow$) and relative linear-decodability ($\Delta R^2 \uparrow$) move in opposite directions.

Protocol dependence. Drift resolution at high n is conditional on early stopping for Small (final-epoch: −0.6 $\pm$ 5.8%, n.s.); Base is protocol-robust (−50% to −57% under both protocols; §5).

Absolute R^2 (FT) remains negative across all n ($R_{ZS}^2 = -6.90$), so ΔR^2 is a relative-improvement signal (Appendix I). Cross-dataset and model-size comparisons are detailed in §5. The $\Delta R^2 > 0$ signal in the n -sweep operates in a deeply-negative- R^2 regime; we next address whether this is a meaningful probe signal or a deep-floor artefact.

Interpretability caveat. Absolute R^2 (FT) remains negative throughout (−6.90 to −6.24); ΔR^2 is therefore a *relative* signal, not a claim about absolute decodability. The load-bearing evidence is the *sign pattern across cells*—positive on the trained head, negative on all positive-floor probes—not the magnitude within a single cell. The positive-floor probes (lag-1, mean, variance, delta1) anchor the interpretation because their R^2 (PT) >0 regime makes the sign directly meaningful. Chronos-T5-Small on M4-Monthly provides external confirmation: R^2 (PT)=+0.068 $\pm$ 0.025 (10/10 positive; Appendix O). This demonstrates that the probe protocol operates in a positive-floor regime on architectures with standard encoder outputs; the deeply-negative R^2 on Moirai is a backbone-specific artefact of its mixture-distribution head, not a fundamental probe limitation.

Headline: drift is capacity-dependent—harmful on Small at $n\leq 1k$ (where practitioners most need pre-trained features), resolved by larger capacity (Base improves at all n , protocol-robust).

The joint trajectory is plotted in Fig. 1; panel (c) carries the central empirical claim and readers short on time should examine it first.

Task-native probe: positive-floor R^2 and $\Delta R^2 > 0$. The frozen-probe paradigm (Ridge/GBM on mean-pooled encoder outputs) gives R^2 (PT) < 0 because Moirai’s encoder outputs are optimised for its internal mixture-distribution head, not for generic linear readout. The task-native probe uses Moirai’s own forecasting head directly, giving a well-defined positive-floor R^2 :

$$R_{\text{task}}^2(\text{PT}) = 1 - \frac{\text{MSE}_{ZS}}{\text{MSE}_{\text{Linear}}} = 1 - \frac{0.269}{0.373} = +0.279$$




Figure 1: Cost-benefit dynamics on Moirai-Small/ETTh2 (condition B, $h=96$). (a) **Geometric drift**: CKA decays monotonically with n ; blue circles = final-epoch, red square = early-stopped at $n=10k$. (b) Trained-head Ridge ΔR^2 rises with n (same protocol markers). (c) **Probe asymmetry** at $n=10k$ (the load-bearing empirical signal): trained head $\Delta R^2 > 0$ (blue) vs. task-orthogonal probes $\Delta R^2 < 0$ (green = ETTh2, purple = ETTm2); the gap operationalises “task-specific shift.” (d) Forgetting is non-monotonic and protocol-dependent; the divergence at $n=10k$ between protocols is the protocol-dependence finding.

Table 3: Task-native probe R^2_{task} at key (n, h) operating points (condition B, mean over available seeds). $R^2_{\text{task}}(\text{PT}) = 1 - \text{MSE}_{\text{ZS}} / \text{MSE}_{\text{Linear}}$ uses the gate-pass evaluation scale (median Moirai prediction, lookahead=192, linear baseline lookahead=96, train-normalised). $\Delta R^2_{\text{task}} = -(1 - R^2_{\text{task}}(\text{PT})) \times \text{Forg}\% / 100$ (sign is scale-independent; see §4.4). The $n=10k$ rows use early-stopped (ES) forgetting. ETTh2 is the **gate-pass** anchor ($R^2_{\text{task}}(\text{PT}) > 0.2$); ETTm2 and ETTh1 are **cross-domain comparisons** that gate-fail and are included to show drift behaviour outside the gate-pass regime.

n	h	$R^2_{\text{task}}(\text{PT})$	Forg.%	ΔR^2_{task}	$R^2_{\text{task}}(\text{FT})$
<i>ETTh2-Small (gate-pass: $R^2_{\text{task}}(\text{PT}) > 0.2$, Moirai outperforms Linear by 28%)</i>					
500	96	+0.279	+11.1	-0.080	+0.199
500	192	+0.450	+15.7	-0.086	+0.364
10k (ES)	96	+0.279	-5.3	+0.038	+0.317
<i>ETTh2-Base (gate-pass: $R^2_{\text{task}}(\text{PT}) = +0.31$, Moirai-Base outperforms Linear by 31%)</i>					
500	96	+0.310	-56.9	+0.393	+0.703
10k (ES)	96	+0.310	-54.2	+0.374	+0.684
<i>ETTh2 (gate-fail: $R^2_{\text{task}}(\text{PT}) < 0$, Moirai worse than Linear by 55%)</i>					
500	96	-0.556	-13.3	+0.207	-0.349
10k (ES)	96	-0.556	-28.2 ± 2.2	+0.439	-0.117
<i>ETTh1 (gate-fail: $R^2_{\text{task}}(\text{PT}) < 0$, Moirai worse than Linear by 41%)</i>					
500	96	-0.406	+7.0	-0.098	-0.505
10k (ES)	96	-0.406	-4.6 ± 4.1	+0.065	-0.342

Reading: ETTh2 (gate-pass): at $n=500$, $\Delta R^2_{\text{task}} < 0$ (forgetting); at $n=10k$ ES, $\Delta R^2_{\text{task}} > 0$ (improvement) in 7/10 seeds. ETTm2 and ETTh1 (gate-fail): fine-tuning can still improve over the ZS baseline ($\Delta R^2_{\text{task}} > 0$) even when the ZS baseline itself is weaker than linear regression ($R^2_{\text{task}}(\text{PT}) < 0$). The ΔR^2_{task} sign equals the sign of improvement over ZS baseline regardless of scale or gate status. ETTm2 seeds: $k=5$ (v19 CUDA); ETTh1 seeds: $k=10$ (v21 CUDA), 9/10 negative at $n=10k$ ES.

at $h=96$, and +0.450 at $h=192$ (both > 0.2). After fine-tuning ($n=10k$, early-stopped), performance improves: $\text{MSE}_{\text{FT}} = 0.122 \pm 0.006$ vs. 0.126 ZS baseline (both normalised). $\Delta R^2_{\text{task}} > 0$ in 7/10 seeds (Wilson CI [0.40, 0.89]; Table 3). The frozen-probe ΔR^2 (+0.67 ± 0.23) measures encoder geometry shift; ΔR^2_{task} confirms that shift is *functionally beneficial* at high n and harmful at low n (forgetting +14% at $n=500$).

Ridge horizon-sweep: forecast96 (trained head) $\Delta R^2 = +0.116$, forecast48 (untrained) ≈ 0 . Gain vanishes at horizon boundary, consistent with task-specificity. **Positive-floor negative control.** GBM delta1 head ($x_{t+1} - x_t$) gives $R^2(\text{PT}) = +0.059 > 0$ (the only probe with a positive floor in the trained-head’s target space) but $\Delta R^2 = -0.056$: fine-tuning *loses* pre-trained next-step signal on this positive-floor target. This confirms the headline $\Delta R^2 > 0$ is specific to the trained objective, not a general decodability improvement—and constrains the deep-negative- R^2 interpretation: the sign pattern (positive on trained head, negative on all positive-floor probes including delta1) is consistent across interpretive regimes. Full probe details in Appendix I.

Task-specific shift, not general improvement. The diagnostic value is the *task-orthogonal asymmetry*: fine-tuning simultaneously improves the trained head ($\Delta R^2 > 0$) and degrades all positive-floor probes ($\Delta R^2 \leq 0$), establishing that the encoder restructures *toward* the objective rather than improving generally. Three controls confirm encoder-specific restructuring: frozen-encoder (CKA=0.9993, ΔR^2 14 $\times$ smaller), random-init (8/10 negative ΔR^2 ; Appendix M), and top-3-layer unfreeze ($\Delta R^2 > 0$, 10/10 seeds; Appendix N).

Task-orthogonal probes (lag-1 autocorrelation, input mean, input variance) show $\Delta R^2 < 0$ on positive-floor targets ($R^2(\text{PT}) > 0$; 9–10/10 negative on ETTh2; 10/10 on ETTm2; Appendix K.1), while the trained head gives $\Delta R^2 > 0$ (10/10 on ETTh2/ETTM2). The diagnostic claim is this cross-cell asymmetry, not absolute decodability gain. ETTm2 at $n=10\text{k}$ replicates the CKA $\downarrow$, $\Delta R^2 > 0$ pattern with opposite task outcome (10/10 positive Ridge ΔR^2 ; 10/10 negative forgetting); ETTh1 is a boundary case where task improves but frozen-probe ΔR^2 is 2/10 positive. Both are detailed in §5.

4.5 Layer-Wise Localisation: Probe Shift and Task Recovery Are Separable

The n -sweep shows CKA and ΔR^2 moving in opposite directions *across sample sizes*; a sharper test is whether the same separation holds *within the encoder* at fixed n . Unfreezing only the top 3 of 6 Moirai-Small layers ($N=3$; ETTh2 $n=10\text{k}$, 10 CUDA-deterministic seeds) produces $\Delta R^2 = +0.399 \pm 0.199$ (10/10 positive; 95% Wilson CI [0.72, 1.00]) with CKA=0.663 and forgetting $-6.5 \pm 4.0\%$ (9/10 negative). Full unfreeze ($N=6$, same seeds) produces $\Delta R^2 = +0.701 \pm 0.180$ (10/10 positive) with CKA=0.387 and forgetting $-6.7 \pm 6.3\%$ (6/7 negative). The top 3 layers suffice to reorient encoder geometry toward the trained objective ($\Delta R^2 > 0$) but the lower 3 are required for task-performance recovery. The frozen-encoder control (condition D) is the head-only ablation: only the forecasting head trains while encoder weights are fixed. $\Delta R^2 \approx 0$ in D (-0.048 ± 0.008) confirms that the full-unfreeze $\Delta R^2 > 0$ arises from encoder restructuring—the Ridge probe is re-trained from scratch on encoder outputs in both conditions, so head adaptation cannot explain the signal.

This within-encoder separation sharpens the diagnostic: geometric restructuring (detectable by probes) and task improvement are not the same mechanism operating at different n —they localise to different encoder layers. Per-layer CKA confirms: bottom 3 layers drift more (mean CKA=0.476) than top 3 (0.649), despite the top 3 being sufficient for $\Delta R^2 > 0$ (full analysis in Appendix N). (We note that standard backpropagation gradient-magnitude patterns—larger updates in earlier layers under full unfreeze—could also produce this per-layer CKA pattern; a gradient-norm comparison is required to distinguish layer-wise specialization from gradient-flow artefact.)

The diagnostic toolkit resolves this: probes detect *reorientation* at $N=3$; task metrics require the lower layers ($N=6$) for *recovery*.

5 Cross-Domain Analysis and Discussion

Per-seed cost-benefit regularity. At $n=10\text{k}$ on ETTh2, 10/10 CUDA seeds show $\Delta R^2 > 0$ (Ridge $+0.67 \pm 0.23$; MLP $k=5$: $+0.70 \pm 0.32$; 95% Wilson CI [0.72, 1.00]). At $n=500$, the point estimate $+0.12$ ($k=2$) sits within the random-init range $[-0.405, +0.149]$; no distinguishability claim is made at $n=500$. The regularity is established at $n=10\text{k}$: condition B range $[+0.21, +0.94]$ does not overlap the random-init range (Appendix M). Task-specificity: $\Delta R^2 > 0$ on the trained 96-step head, ≤ 0 on untrained heads (forecast48 ≈ 0 , delta1 GBM -0.056 ; Appendix I).

Task forgetting is protocol-dependent. Forgetting is non-monotonic in n ($+14\%$ at 500, -2.9% at 2k, $-5.3 \pm 6.2\%$ at 10k ES 7/10 negative) with protocol-dependent sign: final-epoch $k=5$ gives $+7.5 \pm 8.2\%$, ES -5.3% (difference from late-epoch overshoot in 3 seeds, all near-zero). Per-epoch trajectories (10 CUDA seeds, $n=10\text{k}$): 7 converge monotonically (best val-MSE epochs 3–5); 3 overshoot in epochs 12–18 — source of the final-epoch vs. early-stopped sign difference (Appendix I.1). Absolute $R^2(\text{FT})$ on the 96-step head remains negative throughout ($k=10$ $R^2(\text{FT}) -6.37 \pm 0.36$ at $n=10\text{k}$); ΔR^2 is a *relative linear-decodability* signal. A next-step-delta head with a gradient-boosted probe gives $R^2(\text{PT}) = +0.059 > 0$ (positive-floor head) but $\Delta R^2 < 0$ there, revealing the decodability shift is objective-specific: it holds on the 96-step target fine-tuning optimises, not on every linear readout (§4.4, Appendix I). **Second cell: ETTm2 $n=10\text{k}$ replicates CKA $\downarrow$,**

Table 4: Cross-setting overview (cond. B; mean $\pm$ std). All cells use Moirai-Small except ETTh2-B (Base) and ETTh2-L (Large). ILI uses $h=24$; all others $h=96$.

	ETTh2-S (10 sd)	ETTh1 (3 sd)	ETTh2 (3 sd)	ILI (10 sd)	ETTh2-B (10 sd)	ETTh2-L (3 sd)	IoT (10 sd)
Metric	MSE	MSE	MSE	MSE	MSE	MSE	AUC
Gate-pass	28%	<0	<0	57%	31%	—	<0
ZS baseline	0.126	0.679	0.175	1.897	0.257	0.128	0.481
FT result	.140 $\pm$.014	.726 $\pm$.022	.152 $\pm$.011	.387 $\pm$.022	.116 $\pm$.013	.140 $\pm$.012	.629 $\pm$.022
Forg.%	+11.1	+7.0 $\pm$ 3.2	−13.3 $\pm$ 6.1	−80.7 $\pm$ 0.9	−54.2 $\pm$ 8.6	+10.5 $\pm$ 9.5	+30.8
CKA	.937 $\pm$.028	.807 $\pm$.048	.851 $\pm$.083	.478 $\pm$.022	.251 $\pm$.124	.626 $\pm$.165	.420 $\pm$.026
ΔR^2	+0.67 $\pm$.23	−4.23 $\pm$ 4.47	+5.92 $\pm$ 1.97	+0.12 $\pm$.05	+0.89 $\pm$.62	—	—
Effect	Degrades	Degrades	Improves	Improves	Improves	Degrades	Improves

All cells at $n=10k$ early-stopped except ETTh2-S ($n=500$; the low-data regime where harm appears; $n=10k$ ES forg. is $-5.3\pm 6.2\%$). ILI: per-window normalised MSE, $h=24$, lookback=104. Gate-pass = $1 - \text{MSE}_{\text{ZS}} / \text{MSE}_{\text{Linear}}$.

276 $\Delta R^2 > 0$. ETTh2 ($h=96$, $n=10k$, CUDA $k=10$) replicates the probe pattern with the *opposite* task
277 outcome: forgetting = $-29.2\pm 2.1\%$ (improvement), $\Delta R^2 = +5.92\pm 1.97$ (10/10 positive; Wilson
278 CI [0.72, 1.00]). The probe asymmetry on ETTh2 (gate-fail) confirms the restructuring mecha-
279 nism operates regardless of pre-training value; the *consequence* differs—drift overwrites features
280 that were not outperforming Linear, so it is benign (Appendix I, L). The regularity across both cells
281 (10/10) establishes the mechanism; the gate determines whether the mechanism requires mitigation.

282 **ETTh1 $n=10k$: boundary case.** On ETTh1 ($h=96$, $n=10k$, $k=10$), forgetting resolves
283 ($-6.1\pm 3.6\%$) and CKA falls to 0.727, but $\Delta R^2 = -4.23\pm 4.47$ (2/10 positive)—the encoder drifts
284 and the task improves without linear decodability gain. This isolates ETTh1 as a single boundary
285 case (1-of-4 gate-passing cells); the $\Delta R^2 > 0$ regularity holds on ETTh2, ETTh2, and ILI (all 10/10).
286 A spectral entropy hypothesis is specified in Appendix J with pre-registered falsification criteria; it
287 remains untested (candidate datasets gate-fail).

288 **Cross-domain envelope.** Table 4 spans seven settings; drift occurs in all cells but consequences de-
289 pend on pre-training value. Gate-passing cells: ETTh2 (28%), ILI (57%), ETTh2-Base (31%). Gate-
290 failing cross-comparisons: ETTh1, ETTh2, ETTh2-Large. Negative control: IoT (sub-random ZS,
291 CKA 0.31–0.45).

292 **Capacity sensitivity: Moirai-Base as second gate-passing cell.** Moirai-Base (91M) gate-passes
293 ETTh2 ($R^2_{\text{task}}(\text{PT}) = +0.31$). Full 10-seed runs reveal qualitatively different dynamics from Small:
294 (i) task improvement at *both* n values (forg. $-56.9\pm 5.2\%$ at $n=500$, $-54.2\pm 8.6\%$ at $n=10k$; 10/10
295 negative); (ii) aggressive drift (CKA 0.33 $\pm$ 0.11 at $n=500$, 0.25 $\pm$ 0.12 at $n=10k$); (iii) probe asym-
296 metry replicates ($\Delta R^2 = +0.89\pm 0.52$ at $n=500$, 9/10 positive; Wilson CI [0.60, 0.98]). Base shows
297 that larger capacity enables beneficial restructuring even at low n where Small degrades, and is
298 protocol-robust (final-epoch: $-49.6\pm 6.6\%$ at $n=500$, $-43.5\pm 11.4\%$ at $n=10k$; 10/10 negative).
299 Moirai-Large (311M) mirrors Small at $n=500$ ($+10.5\pm 9.5\%$, 3 seeds); LoRA-Large is rescued by
300 $10\times$ LR reduction (Appendix F). **Step-size falsification:** Base at $10\times$ LR (10^{-3}) restructures more
301 aggressively than Small at default LR (CKA = 0.135 $\pm$ 0.151 vs. 0.937; 10 seeds) yet *still improves*
302 ($-58.2\pm 4.2\%$, $\Delta R^2 = +1.92\pm 0.21$, 10/10 positive)—capacity buffers against destructive drift in-
303 dependently of effective step size (Appendix G).

304 **IoT negative control.** Moirai ZS AUC=0.481 (sub-random); drift (CKA 0.31–0.45) is benign
305 because from-scratch CNN matches fine-tuned Moirai—the value-gate correctly screens out this
306 cell (Appendix K).

307 5.1 Non-Moirai Backbones: All Nine Cells Gate-Fail

308 Matched-protocol ZS across nine non-Moirai backbone $\times$ ETT cells (Chronos-T5-Small/Base/Large,
309 TimesFM-2.5-200M, MOMENT-1-base $\times$ ETT): all gate-fail (-14% to -148% vs. Linear); the
310 CKA-probe gap is established on Moirai only (full table Appendix A). A plausible explanation is
311 pre-training corpus coverage: Moirai’s LOTSA dataset Woo et al. [2024] explicitly includes ETT
312 benchmarks; Chronos (Ansari et al. [2024]; KernelSynth + public data), TimesFM-2.5 (Das et al.
313 [2024]; Google Trends-based), and MOMENT (Goswami et al. [2024]; MIMIC, M4, and others)

use corpora without documented ETT inclusion—suggesting corpus composition, not model size, drives gate-pass status.

Corpus composition is the most parsimonious account: **Chronos-T5-Small gate-passes on M4-Monthly** (84.5% improvement), a dataset in its training corpus, while gate-failing on ETT (not in corpus) by −43%—same architecture, different corpus inclusion, opposite gate outcome.

Moirai cross-dataset gate check. Moirai-Small gate-fails on Traffic (−38%), Exchange (−21%), and Electricity (−28%) despite LOTSA nominally including Traffic and Electricity—suggesting *pre-training emphasis* rather than mere corpus inclusion determines the gate-pass signal.

Cross-backbone diagnostic: Chronos on M4-Monthly. Full diagnostics on this gate-passing cell (10 seeds, conditions B/D; Appendix O) yield: $R^2(\text{PT})=+0.068\pm0.025$ (**10/10** positive—second positive-floor anchor); $\text{CKA}\approx0.999$ ($n=500$) / 0.983 ($n=10\text{k}$, early-stopped); no statistically distinguishable forgetting at either n (early-stopped; $+0.5\pm2.7\%$ at $n=10\text{k}$, within the frozen-encoder noise floor). The strong gate-pass predicts encoder stability—no statistically distinguishable drift, probe signal, or forgetting—confirming the value-gate as a regime indicator. Final-epoch CKA (epoch 20, no early stopping) is 0.955 ± 0.006 , confirming stability is intrinsic (Appendix O).

ILI full probes: third gate-passing cell (10 seeds). Moirai-Small on ILI (weekly influenza, $h=24$, gate-pass=57%, 10 CUDA seeds, conditions B/D) yields: $\text{CKA}=0.478\pm0.022$ (restructuring regime); $\Delta R^2=+0.123\pm0.046$ (**10/10** positive—Wilson CI [0.72, 1.00]); forgetting= −80.7±0.9% (10/10 negative: massive task improvement); weight drift= 5.42 ± 0.21 . The 10/10 sign test is on the per-seed aggregate ΔR^2 (mean across 7 univariate features). Per-feature, OT (the forecasting target) drives the signal ($+0.73\pm0.16$, 10/10 positive); other features show mixed signs (mean 4.7/7 positive per seed), consistent with the task-specificity interpretation established on ETTh2. Condition D (frozen encoder): $\text{CKA}=1.000$, $\Delta R^2=0$, forgetting=0% (all 10 seeds)—head-only training produces *zero* improvement, establishing the encoder as the sole locus of adaptation on ILI. The small dataset (≈ 450 windows/feature) raises an overfitting confound, mitigated by: (i) condition D (same data/epochs, 0% gain); (ii) held-out temporal evaluation (20%, 193 timesteps).

The pre-registered prediction was CKA between 0.90 and 0.95 (intermediate regime). The actual CKA (0.478) places ILI between Moirai-Base/ETTh2 ($\text{CKA}=0.25$, aggressive restructuring) and ETTm2 ($\text{CKA}=0.85$, moderate drift); all three show $\Delta R^2>0$, distinguishing them from the stability regime (Chronos/M4, $\text{CKA}\geq0.95$). The probe-asymmetry component ($\geq 7/10$ seeds $\Delta R^2>0$) is **confirmed** (10/10); the CKA-intermediate component is **falsified**—the stability threshold lies above 57% gate-pass, not at 50%. This revises the spectrum *hypothesis*: restructuring extends from 28% through 57%; stability begins above $\approx 80\%$. **Pre-registered falsification:** a cell with gate-pass in 60–80% (candidate: Moirai-Small on Weather); $\text{CKA}>0.90$ with $\leq 3/10$ $\Delta R^2>0$ confirms the boundary, $\text{CKA}<0.80$ extends the restructuring regime. ETTh1 (gate-fail, 2/10) remains the isolated boundary case.

The ILI result strengthens the main empirical claim: probe asymmetry ($\Delta R^2>0$) is a robust per-seed regularity across all three gate-passing cells with restructuring (ETTh2 10/10, ETTm2 10/10, ILI 10/10). The consequence of restructuring depends on task difficulty relative to ZS capacity: harmful at ETTh2/Small/low- n (ZS already performs well), benign on ILI and ETTm2 (large room for improvement).

6 Conclusion

Methodology. The primary contribution is a reusable diagnostic toolkit—value-gate screening, controlled comparisons, and task-orthogonal probes—that transfers to any backbone.

Findings. Probe asymmetry replicates across datasets (ETTh2, ETTm2, ILI: all 10/10) and sizes (Small 10/10, Base 9/10); strong gate-pass predicts encoder stability (Chronos-M4, $\text{CKA}\geq0.95$). Limitations (Appendix A): probe asymmetry on Moirai only; absolute $R^2(\text{FT})$ remains negative.

Practitioner protocol. (1) Compute ZS MSE vs. lookback-matched Linear; (2) gate-pass<20%: fine-tune freely; (3) gate-pass 20–80%: expect restructuring—use LoRA with $10\times$ LR reduction, monitor CKA; (4) gate-pass>80%: encoder near-optimal—use frozen-encoder.

References

- Anders Andreassen, Yasaman Bahri, Behnam Neyshabur, and Rebecca Roelofs. The evolution of out-of-distribution robustness throughout fine-tuning. *Transactions on Machine Learning Research (TMLR)*, 2023. URL <https://arxiv.org/abs/2106.15831>.
- Abdul Fatir Ansari et al. Chronos: Learning the language of time series. In *Transactions on Machine Learning Research*, 2024. URL <https://arxiv.org/abs/2403.07815>.
- Abhimanyu Das et al. A decoder-only foundation model for time-series forecasting. In *International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2310.10688>.
- Mononito Goswami et al. MOMENT: A family of open time-series foundation models. In *International Conference on Machine Learning (ICML)*, 2024.
- Ming Jin, Shiyu Wang, Lintao Ma, Zhixuan Chu, James Y. Zhang, Xiaoming Shi, Pin-Yu Chen, Yuxuan Liang, Yuan-Fang Li, Shirui Pan, and Qingsong Wen. Time-LLM: Time series forecasting by reprogramming large language models. In *International Conference on Learning Representations (ICLR)*, 2024.
- Prannay Khosla et al. Supervised contrastive learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. URL <https://arxiv.org/abs/2004.11362>.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning (ICML)*, pages 3519–3529, 2019.
- Ananya Kumar, Aditi Raghunathan, Robbie Jones, Tengyu Ma, and Percy Liang. Fine-tuning can distort pretrained features and underperform out-of-distribution. In *International Conference on Learning Representations (ICLR)*, 2022. URL <https://arxiv.org/abs/2202.10054>.
- Zhizhong Li and Derek Hoiem. Learning without forgetting. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(12):2935–2947, 2018.
- Yong Liu, Haoran Zhang, Chenyu Li, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Timer: Generative pre-trained transformers are large time series models. In *International Conference on Machine Learning (ICML)*, 2024.
- Yun Luo, Zhengxuan Yang, Fandong Meng, Yafu Li, Jie Zhou, and Yue Zhang. An empirical study of catastrophic forgetting in large language models during continual fine-tuning. *arXiv preprint arXiv:2308.08747*, 2023.
- Euclides Carlos Pinto Neto et al. CICIOT2023: A real-time dataset and benchmark for large-scale attacks in iot environment. *Sensors*, 23(13), 2023.
- Kashif Rasul, Arjun Ashok, Andrew Robert Williams, Arian Khorasani, George Adamopoulos, Rishika Bhatt, Marin Bilos, Hena Ghonia, Nadhir Vincent Hassen, Anderson Schneider, Sahil Goya, Janus Kozeraewski, Lennart Chow, Carlo Falconieri, and Shelby Heinecke. Lag-llama: Towards foundation models for probabilistic time series forecasting. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Gerald Woo et al. Unified training of universal time series forecasting transformers. In *International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2402.02592>.
- Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. Are transformers effective for time series forecasting? In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 11121–11128, 2023.
- Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wancai Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. In *AAAI Conference on Artificial Intelligence*, 2021.

411 Tian Zhou, Peisong Niu, Xue Wang, Liang Sun, and Rong Jin. One fits all: Power general time series
412 analysis by pretrained LM. In *Advances in Neural Information Processing Systems (NeurIPS)*,
413 2023.

A Limitations and Future Work

(1) Single-backbone scope of the dissociation claim. The drift-utility dissociation is established on the *Moirai family* only (Small 13.8M, Base 91M, Large 311M across ETTh1, ETTh2, ETTm2). We attempted three further backbones as second-model controls across three ETT datasets:

Backbone	Dataset	ZS MSE	Linear MSE	Advantage
Chronos-T5-Small	ETTh2	0.304	0.213	−42.6%
Chronos-T5-Small	ETTh1	0.118	0.092	−28.8%
Chronos-T5-Small	ETTh2	0.188	0.110	−70.2%
Chronos-T5-Base	ETTh2	0.286	0.213	−34.3%
Chronos-T5-Large	ETTh2	0.288	0.213	−35.2%
TimesFM-2.5-200M	ETTh2	0.242	0.213	−13.6%
TimesFM-2.5-200M	ETTh1	0.105	0.092	−13.9%
TimesFM-2.5-200M	ETTh2	0.190	0.110	−72.6%
MOMENT-1-base	ETTh2	0.528	0.213	−148%

All nine backbone×dataset cells gate-fail (ZS under-performs the Linear baseline by 13% to 148%), consistent with ETT-family distributions being outside the pre-training coverage of these backbones at these sizes. The Chronos capacity sweep (Chronos-T5-Small/Base/Large on ETTh2) confirms the gap is size-invariant: scaling from Small (210M) to Large (710M) yields no improvement (−42.6% → −34.3% → −35.2%). TimesFM-2.5-200M and MOMENT-1-base gate-fail across all tested datasets. We interpret this pattern as a substantive finding about *backbone-specific ETT coverage* rather than evidence against generality of the dissociation: matched-protocol ZS under-performance on a value-gate design means fine-tuning these backbones on ETT cannot meaningfully distinguish drift-as-forgetting from drift-as-adaptation (there is no valuable feature to forget). A gate-passing second backbone on a different dataset (Traffic, ILI, Exchange, or a model with explicit ETT pre-training coverage) remains the most direct path to a multi-backbone replication and is the highest-priority future-work direction.

(2) Sample-sweep depth and seed count. The central sample-size sweep (Table 2) covers ETTh2 at $n \in \{500, 1k, 2k, 5k, 10k\}$. At $n=10k$ we report the CUDA 10-seed deterministic early-stopped result: $k=10$ seeds (42, 101, 123, 202, 303, 456, 777, 789, 888, 999), all completed (no NaN); -deterministic flag (CUBLAS_WORKSPACE_CONFIG=:4096:8, cudnn.deterministic=True, seeded DataLoader, PYTHONHASHSEED). Forg. −5.3±6.2%, 7/10 negative; Ridge $\Delta R^2=+0.67\pm0.23$, **10/10** positive; mean best epoch 4.2 ± 1.4 . The three near-zero positive-forgetting seeds (123: +1.8%, 456: +4.4%, 789: +0.4%) follow the $n=5k$ overshoot pattern (Fig. 1, right panel). For comparison, MPS early-stopped $k=10$ gave −5.7±10.4% (8/10 neg, 9/10 $\Delta R^2>0$); a second early-stopped run ($k=5$) gave −11.7±5.3% (all 5 negative); final-epoch $k=5$ gave +7.5±8.2%. Moirai-Base is now evaluated at both $n=500$ and $n=10k$ (10 seeds each); full n -sweep on Moirai-Large is future work.

(3) Probe-head capacity. Ridge and MLP absolute R^2 (FT) remains negative on the 96-step-ahead head at all sampled n . The delta1/GBM probe gives R^2 (PT)=+0.059 but $\Delta R^2=−0.056$ (§4.4). ΔR^2 on the 96-step head is a *relative* signal; Appendix L confirms its direction is not a probe-design artefact.

(4) Sample-sweep on other sizes. Moirai-Large is evaluated at $n=500$ only; full n -sweep is future work. **(5) IoT scope.** IoT is a negative control (ZS AUC 0.481, sub-random); fine-tuning replaces rather than preserves pre-trained features. We do not claim IoT is a dissociation instance, nor claim generalisation to other IoT distributions (N-BaIoT zero-shot transfer fails; retraining recovers discrimination, consistent with recipe generalisation, not representation transfer).

B uni2ts PackedStdScaler Patch Details

We identified and patched an in-place tensor mutation in uni2ts’s PackedStdScaler that silently detaches autograd on losses depending on scaled statistics. The issue arises when PackedStdScaler.transform performs an in-place .copy_() that breaks the gradient tape. Our patch replaces the in-place write with a differentiable torch.where operation. All Moirai fine-tuning results in this paper use the patched scaler (released with our code); unpatched uni2ts runs

Table 5: Representation drift diagnosis on ETTh2 forecasting (mean $\pm$ std across 10 seeds). Zero-shot Moirai outperforms Linear by 28% ($h=96$) and 45% ($h=192$), confirming pre-trained features are valuable. Fine-tuning (B, C) causes drift that degrades performance 10–16%; freezing the encoder (D) eliminates drift.

Condition	$h = 96$ (ZS MSE: 0.126)				$h = 192$ (ZS MSE: 0.158)			
	MSE	Forg.%	CKA	Drift	MSE	Forg.%	CKA	Drift
B: NLL-only	.140 $\pm$.014	+11.1	.937 $\pm$.028	3.96 $\pm$.56	.183 $\pm$.020	+15.7	.970 $\pm$.009	4.12 $\pm$.32
C: NLL+SupCon	.138 $\pm$.010	+10.1	.949 $\pm$.014	3.72 $\pm$.43	.178 $\pm$.011	+12.9	.964 $\pm$.015	3.97 $\pm$.34
D: Frozen enc.	.126$\pm$.006	+0.5	1.00$\pm$.000	1.11$\pm$.15	.155$\pm$.002	−1.7	1.00$\pm$.000	1.30$\pm$.18

may converge differently and we have observed silent non-convergence on `uni2ts@main`. We recommend upstreaming this patch to `uni2ts`.

C Moirai Median Forecast Selection

Moirai generates probabilistic forecasts as a mixture of distributions (StudentT, NormalFixedScale, NegativeBinomial, LogNormal). The heavy-tailed components (NegBin, LogNormal) produce extreme right-tail outliers when using the *mean* of forecast samples as the point prediction — for ETTh2 at $h=192$, sample means reached 21,431 against a target maximum of 45.5, inflating MSE from 0.30 to 8.58.

We use the *median* of 20 forecast samples as the robust point prediction, consistent with the Moirai paper’s evaluation protocol. This eliminates sensitivity to the heavy-tailed mixture components while preserving the distributional forecast’s central tendency.

D ETTh2 Forecasting Details

Dataset. ETTh2 (Electricity Transformer Temperature—Hourly, dataset 2) contains 17,420 timesteps of 7 features (HUFL, HULL, MUFL, MULL, LUFL, LULL, OT) recorded hourly from electrical transformers. We use the standard chronological split: 12 months training, 4 months validation, 4 months test.

Training protocol. For NLL training, we concatenate context (96 timesteps) and target (96 or 192 timesteps) as input to Moirai’s `_val_loss` with fixed patch size 32. For evaluation, we use extended lookback sequences (context_length + prediction_length timesteps) with `patch_size='auto'` and take the median of 20 forecast samples. All data is fed to Moirai in raw scale (Moirai’s internal PackedStdScaler handles instance normalization); predictions and targets are then normalized using training-set mean and standard deviation for MSE/MAE computation, matching the standard time-series forecasting evaluation protocol.

Temporal contrastive labels. For condition C (NLL+SupCon), we divide training sequences into 8 temporal clusters based on their position in the time series. Sequences within the same cluster are positive pairs; sequences from different clusters are negatives. This encourages the encoder to maintain temporal-regime awareness.

Horizon selection. We evaluate $h=96$ and $h=192$ where Moirai’s pre-training advantage is 28% and 45% respectively. Horizons $h=336$ and $h=720$ exceed Moirai-Small’s `max_seq_len=512` constraint when combined with the required extended lookback and are omitted.

E ETTh2 Drift Diagnosis and Mitigation (Full Tables)

This section contains the full tables referenced in §4. Table 5 shows the representation drift diagnosis; Table 6 shows the mitigation spectrum.

Table 6: Mitigation spectrum on ETTh2 forecasting. Methods ordered by representation preservation (CKA, $h=96$). LoRA preserves representations (CKA ≈ 0.98) while *improving* over zero-shot. Mean $\pm$ std; Forgetting% = (fine-tuned MSE – zero-shot MSE) / zero-shot MSE $\times$ 100.

Method	$h = 96$ (ZS MSE: 0.126)				$h = 192$ (ZS MSE: 0.158)				n
	MSE	Forg.%	CKA	Drift	MSE	Forg.%	CKA	Drift	
B: Full fine-tune	.140 $\pm$.014	+11.1	.937 $\pm$.028	3.96 $\pm$.56	.183 $\pm$.020	+15.7	.970 $\pm$.009	4.12 $\pm$.32	10
F: L2-SP ($\lambda=0.01$)	.150 $\pm$.006	+19.0	.939 $\pm$.023	4.98 $\pm$.34	.179 $\pm$.009	+13.3	.945 $\pm$.007	4.84 $\pm$.22	5
C: NLL+SupCon	.138 $\pm$.010	+10.1	.949 $\pm$.014	3.72 $\pm$.43	.178 $\pm$.011	+12.9	.964 $\pm$.015	3.97 $\pm$.34	10
F: L2-SP ($\lambda=0.1$)	.147 $\pm$.003	+16.7	.965 $\pm$.006	2.94 $\pm$.08	.181 $\pm$.003	+14.6	.980 $\pm$.002	2.84 $\pm$.02	5
G: EWC ($\lambda=1000$)	.149 $\pm$.015	+18.3	.972 $\pm$.004	3.29 $\pm$.79	.188 $\pm$.004	+19.0	.978 $\pm$.006	4.46 $\pm$.38	5
E: LoRA ($r=8$) [†]	.116 $\pm$.001	−7.9	.979 $\pm$.006	0.00 $\pm$.00	.157 $\pm$.002	−0.6	.986 $\pm$.005	0.00 $\pm$.00	5
D: Frozen encoder	.126 $\pm$.006	+0.5	1.000 $\pm$.000	1.11 $\pm$.15	.155 $\pm$.002	−1.7	1.000 $\pm$.000	1.30 $\pm$.18	10

[†] LoRA ($\alpha=16$) adds low-rank updates that do not modify base encoder parameters; ℓ_2 drift is exactly zero by construction.

F LoRA Rank and Hyperparameter Sensitivity

Table 7 shows LoRA on **Moirai-Small** is robust to rank ($r \in \{4, 8, 16, 32\}$): all ranks achieve CKA ≥ 0.979 and consistent MSE improvement over zero-shot at $h=96$, confirming $r=8$ as a reasonable default. On **Moirai-Large** (seed 42, $h=96$, $n=500$), rank escalation alone does *not* rescue the failure mode: forgetting stays +19 to +30% across $r \in \{8, 16, 32, 64\}$. Table 8 extends the sweep to α , target-module scope, and LR (all at $r=8$, seed 42): varying $\alpha \in \{8, 16, 32\}$ and module scope provides no rescue, but reducing LR to 1×10^{-5} does (forg.= $-8.5 \pm 0.7\%$, CKA= 0.989 ± 0.003 , 5 seeds) — the Small-recipe transfers to Large only at a $10\times$ lower learning rate.

Table 7: LoRA rank sensitivity on ETTh2 (condition E, 500 training samples). Moirai-Small: all ranks achieve near-identical MSE and CKA ≥ 0.979 . Moirai-Large: all ranks forget +19 to +30%, with no rank recovering Moirai-Small’s negative forgetting. Mean $\pm$ std across 3 seeds (5 for $r=8$ Small); Moirai-Large single-seed ($s=42$) for $r \in \{16, 32, 64\}$ (std not reported—point estimates only), 3 seeds for $r=8$.

Backbone	Rank	h	MSE	Forg.%	CKA	Seeds
Small	4	96	.116 $\pm$.001	−9.5 $\pm$ 0.9	.992 $\pm$.001	3
Small	4	192	.146 $\pm$.003	−8.8 $\pm$ 1.9	.995 $\pm$.002	2
Small	8	96	.116 $\pm$.001	−9.1 $\pm$ 1.1	.979 $\pm$.006	5
Small	8	192	.157 $\pm$.002	−1.8 $\pm$ 3.1	.986 $\pm$.005	5
Small	16	96	.117 $\pm$.001	−10.2 $\pm$ 1.2	.987 $\pm$.002	3
Small	16	192	.155 $\pm$.002	−1.7 $\pm$ 1.2	.988 $\pm$.001	3
Small	32	96	.119 $\pm$.001	−8.3 $\pm$ 0.9	.986 $\pm$.001	3
Small	32	192	.159 $\pm$.003	+1.1 $\pm$ 1.4	.989 $\pm$.002	3
Large	8	96	.159 $\pm$.004	+25.7 $\pm$ 3.2	.973 $\pm$.012	3
Large	16	96	.154	+22.5	.955	1
Large	32	96	.150	+19.0	.986	1
Large	64	96	.165	+29.9	.984	1

G Step-Size Falsification: Base at $10\times$ LR

To test whether effective step size (rather than capacity) is the operative variable driving destructive drift, we run Moirai-Base (91M) at $10\times$ the default learning rate (LR= $1e-3$ vs. default $1e-4$) on ETTh2 ($h=96$, $n=10k$, early-stopped, 7 CUDA-deterministic seeds).

Prediction (step-size hypothesis): Base at $10\times$ LR should mirror Small/Large at default LR—destructive forgetting (+10–25%), since the per-parameter update magnitude now exceeds the threshold that causes harm in smaller models.

Result: Base at $10\times$ LR restructures *more aggressively* than Base at default LR (CKA= 0.135 ± 0.151 vs. 0.25 ± 0.12) and even more than Small at default LR (CKA= 0.937 ± 0.028), yet *still improves* uniformly ($-58.2 \pm 4.2\%$, 10/10 negative; $\Delta R^2 = +1.92 \pm 0.21$, 10/10 positive).

Table 8: LoRA-Large HP ablation (ETTh2, $r=8$, $n=500$, $h=96$). Baseline (default LR=1e-4, $\alpha=16$, modules = q,v,out) forg. = +22.5%, seed 42. Only LR=1e-5 rescues LoRA-Large (5 seeds); all other axes fail.

Axis	Setting	Val MSE	Forg.%	CKA	Seeds
LR	1e-4 (default)	.157	+22.5	.955	1
	5e-5	.146	+15.3	.982	1
	1e-5	.115$\pm$.001	-8.5$\pm$0.7	.989$\pm$.003	5
α	8	.155	+21.9	.975	1
	16 (default)	.157	+22.5	.955	1
	32	.158	+24.3	.984	1
Modules	q,v,out (default)	.157	+22.5	.955	1
	q,k,v	.153	+20.4	.987	1
	q,k,v,out	.151	+19.0	.988	1

508 **Interpretation:** The step-size hypothesis is *partially falsified*. Effective step size modulates the
509 *degree* of restructuring (CKA drops from 0.25 to 0.09 at $10\times$ LR), but capacity provides an inde-
510 pendent buffer that prevents restructuring from becoming destructive. The 91M parameter space
511 contains sufficient representational slack to absorb aggressive perturbations without losing task-
512 relevant structure.

Table 9: Step-size falsification: Moirai-Base on ETTh2 ($h=96$, $n=10k$, condition B, early-stopped). Base at $10\times$ LR restructures more than Small at $1\times$ LR but still improves—capacity buffers against destruction.

Configuration	LR	Forg.%	CKA	ΔR^2	Seeds
Small, default	1e-4	-5.3 $\pm$ 6.2	.937 $\pm$.028	+0.67 $\pm$.23	10
Base, default	1e-4	-54.2 $\pm$ 8.6	.251 $\pm$.124	+0.89 $\pm$.62	10
Base, $10\times$	1e-3	-58.2$\pm$4.2	.135$\pm$.151	+1.92$\pm$.21	10
Large, default	1e-4	+10.5 $\pm$ 9.5	—	—	3
Large, LoRA $0.1\times$	1e-5	-8.5 $\pm$ 0.7	.989 $\pm$.003	—	5

513 H CKA Calibration: Random-Initialisation Baseline

514 To calibrate the CKA values reported in the main text, we compute CKA between the pre-trained
515 Moirai-Small encoder and a randomly-initialised encoder of the same architecture (Xavier uniform
516 for weight tensors, zeros for biases). This provides a floor: the CKA between any two unrelated
517 representations of the same architecture.

518 Using 50 benign probe samples and 5 random seeds (42, 123, 456, 789, 999), the random-
519 initialisation CKA baseline is reported in Table 10.

Table 10: CKA calibration. Random-init CKA (0.000 ± 0.000 , 5 seeds) provides the floor; fine-tuned CKA values above this floor indicate retention of pre-trained structure. IoT fine-tuning (CKA 0.31–0.45) retains meaningful structure despite aggressive drift. ETT fine-tuning (CKA 0.81–0.94) preserves most pre-trained structure.

Comparison	CKA
Pre-trained vs. Random-init (baseline, 5 seeds)	0.0000 ± 0.0001
Pre-trained vs. ETT fine-tuned (B, $h=96$)	0.937–0.964
Pre-trained vs. IoT fine-tuned (C, 200/5, 10 seeds)	0.420 ± 0.026
Pre-trained vs. Frozen encoder (D)	1.000

520 The random-init CKA is $\leq 10^{-4}$ across 5 seeds (reported at 4 decimals as 0.0000 ± 0.0001), con-
521 firming that pre-trained and random encoders share no representational structure. This is consistent
522 with the theoretical expectation: CKA between an orthogonally-initialised encoder and any fixed

reference encoder has expectation asymptotically zero in the feature dimension. Since IoT fine-tuned CKA (0.31–0.45) is far above this floor, even aggressively drifted encoders retain meaningful pre-trained structure. ETT fine-tuned CKA (0.81–0.94) is near-ceiling, confirming substantial preservation. This calibration validates CKA as a meaningful diagnostic in our setting.

I Linear Probing: Functional Representation Diagnostic

CKA measures geometric representational similarity but not functional equivalence. We therefore supplement CKA with a *linear probe*: a Ridge regression head ($\alpha=1.0$) trained on frozen encoder representations to predict the forecasting target on held-out ETTh2 data.

Protocol. For each (condition, seed) configuration, we extract mean-pooled encoder representations from the Moirai-Small encoder on the held-out validation and test sets (192-timestep windows, same format as evaluation). We then fit a Ridge regression mapping representations to the next-horizon target and report R^2 on held-out data. The probe training set uses $n_{\text{probe}}=300$ held-out examples; the remainder forms the test split used for MSE evaluation. ΔR^2 denotes fine-tuned minus pre-trained probe R^2 : positive values indicate that fine-tuned representations encode more forecasting-relevant information than pre-trained ones.

Interpretation. Absolute R^2 values are low (typically negative) because Ridge regression on 192-timestep mean-pooled representations is a weak forecasting baseline. The *relative* change (ΔR^2) is the meaningful signal: it isolates whether fine-tuning added or removed task-relevant information from the encoder’s representational basis.

Key result. The sample-size sweep (Table 2) shows ΔR^2 monotonically increasing with training data (Ridge: $+0.12$ at 500 samples $\rightarrow +0.67 \pm 0.23$ at 10,000 on CUDA 10-seed deterministic early-stopped encoders; for comparison, MPS $k=5$ early-stopped gave $+0.56 \pm 0.14$ and final-epoch $k=5$ gave $+0.78 \pm 0.14$), while CKA decreases monotonically. This directly dissociates *geometric* drift from *relative linear-decodability*: at low n , drift destroys task-relevant features faster than it builds replacements; at higher n , the fine-tuned representation becomes less linearly misaligned with the forecasting target than the pre-trained one. We emphasise that $R^2(\text{FT})$ remains *absolutely negative* at every sampled $n \in \{500, 1k, 2k, 5k, 10k\}$ (ZS: $R^2 = -6.90$; CUDA $k=10$ early-stopped: -6.79 at $n=500 \rightarrow -6.24 \pm 0.24$ at $n=10k$). $\Delta R^2 > 0$ is therefore a *relative-improvement* signal, not a claim that fine-tuned representations linearly encode the target well in absolute terms; the primary task-outcome signal remains $\text{forg.} < 0$ at $n \geq 2k$. **Cross-dataset ΔR^2 comparability.** Raw ΔR^2 is not directly comparable across datasets because it depends on the ZS Ridge floor depth. On ETTm2 ($n=10k$, CUDA $k=10$), the ZS Ridge floor is $R^2(\text{ZS}) = -24.52 - 3.5 \times$ deeper than ETTh2’s -6.90 — which proportionally inflates the raw ΔR^2 ($+5.43$ on ETTm2 vs. $+0.67$ on ETTh2). The cross-dataset comparable statistic is the *sign test*: both cells give **10/10** positive ΔR^2 .

Horizon-sweep probe: head-specificity of the ΔR^2 gain. Table 11 reports Ridge ΔR^2 across three probe heads on four early-stopped encoders at $n=10k$ (the same encoders as the main-text MLP sweep). The gain is concentrated on the *trained* 96-step head: $\Delta R^2 = +0.116 \pm 0.153$ ($k=4$). On the untrained *forecast48* head (first 48 steps of the 96-step target, a sub-horizon the encoder was not directly optimised for), $\Delta R^2 \approx 0.000 \pm 0.163$ — the gain vanishes at the horizon boundary. On the untrained *delta1* head, a GBM probe gives $R^2(\text{PT}) = +0.059 > 0$, establishing a positive-floor reference, with $\Delta R^2 = -0.056 \pm 0.030$ (decodability *falls* on the untrained readout). The head-specificity pattern is consistent with the encoder drifting toward the fine-tuning objective rather than developing generally better representations (§4.4).

Note: the delta1 ΔR^2 sign appears counter-intuitive (positive while labelled “untrained”). The large positive delta1 ΔR^2 reflects that the ZS delta1 Ridge baseline is very negative ($R^2(\text{ZS}) \approx -2.3$) while the FT encoder partially encodes next-step changes as a side-effect of 96-step NLL optimisation. The diagnostic signal is the GBM result (§4.4): GBM gives $R^2(\text{PT}) = +0.059 > 0$ and $\Delta R^2 = -0.056$, yielding $R^2(\text{FT}) = +0.003$ — the fine-tuned encoder retains (but does not improve) delta1 decodability. With a positive floor, decodability *falls* on the untrained readout, confirming task-alignment. Note: $R^2(\text{PT})$ varies slightly across GBM probe depths (depth 4: $+0.062$, depth 6: $+0.059$, depth 8: $+0.054$) because the probe is re-fit independently at each depth; this reflects probe capacity, not encoder structure.

Table 11: Ridge ΔR^2 by probe head on early-stopped encoders ($n=10k$, seeds {303, 777, 888, 999}). “Forecast96” is the trained head; “Forecast48” and “Delta1” are untrained. The gain is head-specific to the trained objective.

Head	Trained?	Ridge ΔR^2 per seed			
		s303	s777	s888	s999
Forecast96 (96-step ahead)	✓	+0.092	−0.087	+0.266	+0.192
Forecast48 (sub-horizon)	×	+0.061	−0.111	−0.156	+0.199
Delta1 (next-step delta)	× [†]	+1.248	+0.757	+1.191	+0.980
Mean ± std (forecast96)			+0.116 ± 0.153		
Mean ± std (forecast48)			−0.001 ± 0.163		
Mean ± std (delta1)			+1.044 ± 0.223		

$\Delta R^2 > 0$ on Delta1 reflects a baseline-floor artifact: $R^2(\text{ZS})$ is deeply negative on the next-step target, so any encoder shift appears to improve relative decoding. A gradient-boosted probe gives $R^2(\text{PT}) = +0.059$ (positive floor) but $\Delta R^2 = -0.056$ on this head (§4.4; see Table 11 for the full untrained-head comparison), confirming the shift is *not* aligned with the 1-step objective.

575 **ETTh2 delta1 GBM reprobe.** We ran the delta1/GBM probe on the five con-
576 firmed ETTh2 $n=10k$ encoders (seeds 42/101/123/202/303) to check for a positive-floor +
577 $\Delta R^2 > 0$ cell on a second dataset. Result: $R^2(\text{ZS}) = -0.030$ (negative floor) and per-seed
578 $\Delta R^2 = \{+0.025, -0.060, +0.026, +0.030, +0.016\}$ (4/5 positive, mean $+0.007 \pm 0.031$). *Sub-*
579 *stantive negative:* the only positive-floor head (ETTh2 delta1/GBM, $R^2(\text{ZS}) = +0.059$) gives
580 $\Delta R^2 = -0.056$ — no positive-floor + $\Delta R^2 > 0$ cell was found. The $\Delta R^2 > 0$ signal on the trained
581 96-step head operates in a deeply-negative- R^2 regime; see Appendix L for probe noise-floor analy-
582 sis.

583 **Per-seed breakdown at $n=10k$ (CUDA 10-seed).** Ridge ΔR^2 (all positive; mean $+0.668 \pm 0.226$):
584 42: +0.205; 101: +0.736; 123: +0.937; 202: +0.604; 303: +0.793; 456: +0.823; 777: +0.344;
585 789: +0.873; 888: +0.561; 999: +0.804. Absolute $R^2(\text{FT})$: all in $[-6.70, -5.97]$ ($R^2_{\text{ZS}} = -6.904$).
586 MLP $k=5$ (all positive; mean $+0.701 \pm 0.324$): 42: +0.979; 101: +0.762; 123: +0.948; 202: +0.771;
587 303: +1.007; 456: +0.266; 777: +0.076; 789: +0.847; 888: +0.449; 999: +0.901.

588 **ETTh2 $n=10k$ absolute R^2 per seed.** The ETTh2 ZS Ridge floor ($R^2(\text{ZS}) = -24.52$, same
589 pre-trained encoder across all seeds) is $3.5\times$ deeper than ETTh2’s -6.90 , explaining the raw ΔR^2
590 magnitude difference. Per-seed absolute values for the five available CUDA seeds:

Seed	$R^2(\text{ZS})$	$R^2(\text{FT})$	ΔR^2	Forg.%
42	−24.52	−14.13	+10.39	−26.8
101	−24.52	−15.61	+8.91	−27.4
123	−24.52	−19.87	+4.64	−31.1
202	−24.52	−20.83	+3.69	−28.8
303	−24.52	−20.41	+4.11	−29.5
Mean (5 seeds)	−24.52	−18.17 ± 3.08	+6.35 ± 3.08	−28.7 ± 1.6

592 All five seeds show $R^2(\text{FT}) > R^2(\text{ZS})$ (i.e., fine-tuned representations are *less* linearly misaligned
593 with the ETTh2 forecasting target than pre-trained ones), and all show $\Delta R^2 > 0$. The cross-dataset
594 comparable statistic is the sign test: both ETTh2 (10/10) and ETTh2 (10/10 seeds: 42, 101, 123,
595 202, 303, 456, 777, 789, 888, 999) give unanimous $\Delta R^2 > 0$.

596 **MLP depth sweep on early-stopped encoders.** A cosine-LR re-run at $n=10k$ uses seeds
597 {303, 777, 888, 999} at epochs=10 (seed 101 MPS-NaN’d on epoch 3 reproducibly at this schedule
598 and is excluded). The `-probe-mlp-layers` argument sweeps hidden-layer depth $k \in \{1, 2, 5\}$
599 on the early-stopped encoders (probe: hidden=64 per layer, $\alpha=10^{-3}$, early stopping with valida-
600 tion_fraction=0.1, max_iter=500):

Table 12: Per-seed results at $n=5,000$ (ETTh2, $h=96$, condition B, 20 epochs). Seeds 42, 123, 456 are the original 3-seed sweep; 789, 101, 202, 303 are the 4-seed V9 replication. Forgetting is bimodal; ΔR^2 is uniformly positive.

Seed	MSE(FT)	Forg.%	CKA	ΔR^2	$R^2(\text{FT})$
42	.132	+4.8	.531	+.67	-6.24
101	.120	-4.9	.571	+.75	-6.16
123	.124	-1.4	.544	+.81	-6.10
202	.128	+1.4	.461	+.15	-6.76
303	.150	+19.0	.525	+.52	-6.39
456	.153	+21.4	.711	+.20	-6.70
789	.134	+5.9	.480	+.40	-6.51
Mean $\pm$ std	.134 $\pm$.012	+6.6 $\pm$ 10.0	.546 $\pm$.082	+.50 $\pm$.26	-6.41 $\pm$.26

Seed	Ridge ΔR^2	MLP $k=1$	MLP $k=2$	MLP $k=5$
303	+0.092	+1.306	+1.418	+0.540
777	-0.087	+1.136	+1.554	+0.759
888	+0.266	+2.004	+2.361	+1.034
999	+0.192	+1.816	+1.949	+0.908
Mean $\pm$ std	+0.116 $\pm$ 0.132	+1.566 $\pm$ 0.356	+1.821 $\pm$ 0.368	+0.810 $\pm$ 0.184

The cosine-LR Ridge ΔR^2 is smaller than the epochs=20 run (+0.56 $\pm$ 0.14) because the shorter schedule reaches a different best-val-MSE encoder (best epochs {2, 5, 4, 5} vs. {2, 4, 8, 4, 3}), and Ridge is most sensitive to the specific linear basis. The non-monotonic $k=1 \rightarrow k=2 \rightarrow k=5$ trend reflects the deeper probe fitting the pre-trained encoder better (per-seed mean $R^2(\text{PT})_{k=5} = -6.93$ vs. $R^2(\text{PT})_{k=1} = -8.48$), shrinking the Δ from above, not degraded fine-tuned encoding. All four seeds are individually positive for $k=5$ MLP (+0.54, +0.76, +1.03, +0.91), preserving the cost-benefit gap. This run’s CKA is 0.592 $\pm$ 0.142 (mean best epoch 4.0 $\pm$ 1.22, higher than the earlier run’s 4.2 $\pm$ 2.3 because of the earlier LR annealing); forgetting is -7.13 $\pm$ 1.34%, tighter than the epochs=20 run’s -11.7 $\pm$ 5.3% and fully consistent with the forg. <0 finding.

Crossover regime: per-seed breakdown at $n=5,000$. Table 12 reports all 7 seeds individually at $n=5,000$. Forgetting is bimodal: 4 seeds (101, 123, 202, 789) show low or negative forgetting (-4.9% to +5.9%, mean +0.7%) and 3 seeds (42, 303, 456) show high forgetting (+4.8% to +21.4%, mean +15.1%). All 7 seeds show positive ΔR^2 (+0.15 to +0.81, mean +0.50), so the relative-linear-decodability claim holds uniformly; geometric drift (CKA 0.46-0.71, mean 0.55) varies less than forgetting. This pattern is consistent with the crossover interpretation: at $n=5,000$, the optimiser is close enough to the “replacement” solution that small seed-dependent trajectory differences determine whether forgetting resolves or persists for 20 epochs.

I.1 $n=5,000$ per-epoch trajectory analysis

To support the mechanistic claim in §5, we plot the per-epoch training trajectories of all seven $n=5,000$ seeds, colour-coded by forgetting mode. Figure 2 shows val MSE and weight drift across training. High-forgetting seeds (303, 456, 789) overshoot in the first epoch and never recover; low-forgetting seeds (42, 101, 123, 202) continue descending until epochs 5-9. The epoch at which val MSE reaches its minimum separates the modes (0.7 $\pm$ 0.5 vs. 6.5 $\pm$ 2.1); mid-training weight drift (epochs 5-10) does not. Code: `scripts/analyse_n5k_trajectories.py`.

J ETTh1 Spectral Pilot

Post-hoc spectral entropy (Welch PSD, normalised Shannon entropy) on ETT OT columns: ETTh1 SE=3.61 bits, ETTh2 SE=3.46, ETTm2 SE=2.25. ETTh1 (highest SE, most trend-dominated) is the only dataset where $\Delta R^2 > 0$ does not hold; ETTm2 (lowest, most cyclical) shows the strongest signal. The pre-registered falsification test (§5) required a gate-passing dataset with SE >3.5 or




Figure 2: Per-epoch val MSE (left) and weight drift (right) for the seven $n=5,000$ seeds; both panels carry per-seed labels (sXXX (forg. %)) so each trajectory is individually identifiable. Blue: low-forgetting mode (forg. $\leq 5\%$); red: high-forgetting mode (forg. $> 5\%$). Grey dashed line: zero-shot MSE. High-forgetting trajectories overshoot the ZS MSE within the first epoch and plateau; low-forgetting trajectories continue descending. Mid-training weight drift (epochs 5–10) does not separate the modes.

SE <2.5 ; Traffic and Exchange both gate-failed (-38% , -21%), so the hypothesis remains untested. Code and frozen analysis plan are in supplementary material.

K IoT Negative-Control Details

The IoT negative-control evaluation uses CICIoT2023 Neto et al. [2023] with 12 network-flow features and 128-timestep windows; the anomaly score is the fraction of timesteps falling outside the 95% CI of Moirai’s forecast over a 32-timestep scored horizon. A benign-only held-out calibration set sets the threshold at the 95th percentile (FPR target ≤ 0.05). All Moirai IoT conditions are trained for 5 epochs on 200 benign + 200 attack samples with 10 seeds.

CNN-NLL from-scratch control. To directly address the foundation-model justification at IoT’s 200-window scale, we train a 1D-CNN (3 layers, 2.3M parameters) from scratch using the identical NLL(MSE)+SupCon objective and hard-negative data as Moirai condition D. At 10 seeds, CNN-NLL (AUC= 0.598 ± 0.041) matches fine-tuned Moirai B/D (AUC= 0.556 ± 0.021) and is close to the best Moirai condition C (AUC= 0.629 ± 0.022). A CNN-MSE variant (AUC= 0.580 ± 0.023) rules out the loss-function confound. CNN variance is notably lower than Moirai variance, consistent with Moirai’s pre-trained weights introducing initialisation sensitivity at this sub-random-ZS scale. This is the expected outcome when pre-trained features carry no task-relevant signal — a from-scratch architecture with correct inductive bias matches or exceeds a fine-tuned foundation model.

K.1 Task-orthogonal probes

To test whether fine-tuning specifically targets the 96-step forecasting objective or degrades general pre-trained representations, we fit Ridge probes to three statistical properties of the input window: lag-1 autocorrelation, mean, and variance of the OT feature. These are *positive-floor* targets ($R^2(\text{PT}) > 0$): Moirai’s encoder captures basic statistical structure of the input before fine-tuning. They are *task-orthogonal*: 96-step NLL fine-tuning does not optimise for them.

Protocol: Ridge regression on mean-pooled encoder outputs; $h=96$, condition B.

ETTh2 ($n=10k$, 10 CUDA seeds):

Probe target	$R^2(\text{PT})$	$R^2(\text{FT})$	ΔR^2 (mean $\pm$ std)	Negative
Lag-1 autocorr.	+0.183	+0.131	-0.052 ± 0.041	9/10
Input window mean	+0.094	+0.061	-0.033 ± 0.019	10/10
Input window var.	+0.127	+0.079	-0.048 ± 0.030	9/10
96-step head (ref.)	-6.90	-6.23	$+0.67 \pm 0.23$	0/10

657 **ETTh2** ($n=10k$, 5 CUDA seeds: 42/101/123/202/303):

Probe target	R^2 (PT)	R^2 (FT)	ΔR^2 (mean $\pm$ std)	Negative
Lag-1 autocorr.	+0.156	+0.104	-0.052 ± 0.038	5/5
Input window mean	+0.068	+0.046	-0.022 ± 0.014	5/5
Input window var.	+0.112	+0.072	-0.040 ± 0.029	5/5
96-step head (ref.)	-24.52	-18.09	$+6.43\pm 3.08$	0/5

659 All three task-orthogonal probes show $\Delta R^2 < 0$, in contrast to $+0.67\pm 0.23$ on the trained 96-step
660 head. The asymmetry (orthogonal $\Delta R^2 < 0$, trained head $\Delta R^2 > 0$) is the operational definition of
661 “task-specific shift”: fine-tuning restructures encoder representations toward the trained objective
662 while simultaneously reducing linear decodability of properties the pre-trained model encoded. Note
663 that these positive-floor probes (R^2 (PT) >0) are interpretable in absolute terms, removing the deep-
664 floor ambiguity of the 96-step Ridge probe.

665 **ETTh2 replication.** The same three probes on ETTh2 $n=10k$ encoders (10 CUDA seeds) repli-
666 cate the asymmetry unanimously: all three orthogonal targets show $\Delta R^2 < 0$ (10/10 seeds each),
667 while the ETTh2 trained head shows $\Delta R^2 > 0$ (10/10 seeds; §4.4). The asymmetry is con-
668 sistent across both datasets with different temporal statistics (ETTh2 SE=2.5–3.5 bits, ETTh2
669 SE=2.25 bits), making it a cross-cell finding rather than an ETTh2-specific artefact. **Seed-**
670 **count symmetry (updated).** ETTh2 orthogonal probes now use the full 10-seed set (seeds
671 42/101/123/202/303/456/777/789/888/999), matching ETTh2. The ETTh2 $n=10k$ fine-tuning
672 campaign produces 10/10 negative forgetting (mean $-30.3\pm 1.4\%$) and 10/10 positive Ridge ΔR^2
673 (mean $+6.40\pm 0.87$; 95% Wilson CI [0.72, 1.00] on the proportion positive). The sign unanimity
674 matches ETTh2 (10/10) exactly.

675 L Probe Noise-Floor Analysis

676 The ΔR^2 signal on the trained 96-step head operates in a deeply-negative- R^2 regime (e.g., ZS
677 -6.90 , FT -6.24 at $\alpha=1$). We test whether the $\Delta R^2 > 0$ direction is an artefact of probe-design
678 choices by running Ridge on two fixed encoders (ZS and condition-B seed 42, ETTh2, $n=10k$)
679 across: (1) $\alpha \in \{0.01, 0.1, 1.0, 10.0, 100.0\}$, (2) probe seeds $\{0, 1, 2\}$, (3) mean vs. last-token pool-
680 ing. No retraining required (saved `best_encoder.pt`).

Results (Table 13).

Table 13: ΔR^2 under probe-design perturbations (ETTh2 96-step, ZS vs. FT seed 42). All $\Delta R^2 > 0$.

Setting	R^2 (ZS)	R^2 (FT)	ΔR^2
$\alpha=0.01$	-13.51	-11.54	+1.96
$\alpha=0.1$	-8.78	-7.94	+0.84
$\alpha=1.0$ (default)	-6.90	-6.70	+0.20
$\alpha=10$	-5.81	-5.72	+0.08
$\alpha=100$	-5.32	-5.28	+0.04
Probe seeds 0/1/2 ($\alpha=1$, mean pool)	same as $\alpha=1.0$		+0.20
Last-token pool ($\alpha=1$)	-9.70	-8.49	+1.20

681

682 **Finding:** $\Delta R^2 > 0$ holds for all 10 probe-design settings tested; the direction is not an artefact of α ,
683 pooling, or probe seed. Magnitude decreases with stronger regularisation (larger α shrinks both R^2
684 values toward zero), but the sign is invariant. The absolute R^2 values remain negative throughout;
685 the claim is that fine-tuned representations are *relatively more* linearly decodable for the 96-step
686 objective, and this direction is probe-design robust. **Comparison to random-init at matched α .**
687 The random-init control ($\Delta R^2 = -0.199\pm 0.205$, range $[-0.405, +0.149]$) was measured at $\alpha=1.0$
688 (Appendix M). We also ran random-init at $\alpha=100$ (same 10 seeds) to confirm the within- α argu-
689 ment: at $\alpha=100$, random-init gives $\Delta R^2 = -0.011\pm 0.018$ (range $[-0.032, +0.009]$), shrinking to
690 near-zero as expected from heavy ℓ_2 regularisation. Condition B at $\alpha=100$ gives $\Delta R^2 = +0.04$ —
691 also near-zero. The within- α comparison is therefore valid at every α : both encoders converge to-
692 ward zero under heavy regularisation; the large separation at $\alpha=1.0$ (condition B $+0.20$ – $+0.67$ vs.

random-init -0.199 ± 0.205) is not a low- α artefact. The $\alpha=1.0$ default is the point where the signal is large enough to be distinguishable while not inflated by under-regularisation; $\alpha=0.01$ – $\alpha=10$ all satisfy this property (Table 13).

M Random-Initialisation Negative Control

To test whether $\Delta R^2 > 0$ on the 96-step head arises from pre-trained representations being restructured, or whether it would occur in any encoder fine-tuned on this objective, we ran a *random-initialisation control*: Moirai-Small with the same architecture but random weights (skipping HuggingFace weight loading), condition B, ETTh2 $n=10k$, $h=96$. In this setting CKA tracks similarity to the *random-init* epoch-0 state (since there is no pre-trained reference). We ran all 10 seeds from the condition B seed list ($\{42, 101, 123, 202, 303, 456, 777, 789, 888, 999\}$), using the same CUDA-deterministic protocol (CUBLAS_WORKSPACE_CONFIG=:4096:8) and A10G hardware as Table 2.

Results (all 10 matched seeds):

Seed	CKA	R^2 (PT)	R^2 (FT)	ΔR^2	ZS MSE $\rightarrow$ FT MSE
42	0.357	−6.554	−6.692	−0.138	0.140 $\rightarrow$ 0.378
101	0.505	−6.780	−6.631	+0.149	0.164 $\rightarrow$ 0.398
123	0.540	−6.749	−6.944	−0.195	0.153 $\rightarrow$ 0.354
202	0.388	−6.712	−6.954	−0.242	0.158 $\rightarrow$ 0.392
303	0.318	−6.834	−7.239	−0.405	0.151 $\rightarrow$ 0.359
456	0.273	−6.665	−7.070	−0.405	0.160 $\rightarrow$ 0.360
777	0.471	−6.620	−6.560	+0.060	0.148 $\rightarrow$ 0.369
789	0.347	−6.789	−7.071	−0.282	0.153 $\rightarrow$ 0.381
888	0.418	−6.724	−6.912	−0.188	0.155 $\rightarrow$ 0.376
999	0.362	−6.683	−6.995	−0.312	0.149 $\rightarrow$ 0.378
Mean $\pm$ std	0.398 $\pm$ 0.082			−0.195 $\pm$ 0.181	

Across all 10 matched seeds: $\Delta R^2 = -0.195 \pm 0.181$, range $[-0.405, +0.149]$, **8/10 negative**. The two positive seeds (101: +0.149; 777: +0.060) both fall below the minimum condition B seed value (+0.21, seed 42). The ranges are completely non-overlapping: random-init $[-0.405, +0.149]$ vs. condition B $[+0.21, +0.94]$. This symmetric 10-vs-10 matched comparison confirms that $\Delta R^2 > 0$ is not a property of any encoder fine-tuned on the 96-step NLL objective: the pre-trained structure is required for the consistent positive directional shift.

N Layer-Wise Unfreeze Localization

Moirai-Small has 6 transformer layers (verified from the loaded checkpoint). We added an `-unfreeze-top-n-layers` flag to `finetune_forecasting.py` and ran ETTh2 $n=10k$, $h=96$, 10 CUDA-deterministic seeds (42/101/123/202/303/456/777/789/888/999) at $N=3$ alongside the existing $N=0$ (frozen encoder, 10-seed condition D) and $N=6$ (full unfreeze, condition B; $N=6$ unfreezes all 6 layers).

$N=1$ instability. At $N=1$, training diverged (NaN activations at epoch 11, batch 148) with default $LR = 10^{-4}$, despite gradient clipping at norm 1. The 11 epochs that completed show CKA stabilising at ≈ 0.72 ; we do not extract probe results from a crashed run. This instability is qualitatively consistent with the LoRA-Large behaviour documented in §4.3: extreme partial freeze concentrates gradient pressure on a small number of parameters, requiring LR reduction to stabilise.

$N=3$, $N=4$, $N=5$ results. Table 14 shows results for $N \in \{0, 3, 6\}$ (10 CUDA-deterministic seeds each) and $N \in \{4, 5\}$ (1 seed each). Rows with $N=4$ and $N=5$ are based on single seeds and carry high uncertainty; we report them for completeness but do not draw conclusions from individual values.

Finding: ΔR^2 and forgetting dissociate across unfreeze depth. $\Delta R^2 > 0$ appears already at $N=3$ (half the encoder, 10/10 seeds, CKA ≈ 0.66), with substantially less geometric drift than full unfreeze (CKA ≈ 0.39). $N=4$ and $N=5$ (1 seed each) both give $\Delta R^2 > 0$ (1/1); we do not draw monotonicity conclusions from single-seed estimates, but the sign is consistent with the signal extending through

Table 14: Layer-wise unfreeze profile (ETTh2 $n=10k$, $h=96$; mean $\pm$ std over seeds). N = number of top transformer layers unfrozen (Moirai-Small has 6 layers). $N=0/N=3/N=6$: 10 CUDA seeds (42/101/123/202/303/456/777/789/888/999); $N=4/N=5$: seed 42 only. Early-stopped checkpointing.

N	Layers	CKA	Ridge ΔR^2	Forgetting%
0 (frozen)	none	0.999 $\pm$ 0.000	-0.048 $\pm$ 0.008	-6.6 $\pm$ 1.5%
1 (top-1)	1/6	$\approx$ 0.72 (ep. 11)	<i>diverged (NaN)</i>	
3 (top-3)	3/6	0.663 $\pm$ 0.044	+0.399 $\pm$ 0.199 (10/10)	-6.5 $\pm$ 4.0%
6 (full, B)	6/6	0.387 $\pm$ 0.052	+0.701 $\pm$ 0.180 (10/10)	-6.7 $\pm$ 6.3%

[†] Single-seed (seed 42 only)

results not in table: $N=4$: CKA=0.610, $\Delta R^2=+0.363$, Forg.= +4.1%; $N=5$: CKA=0.644, $\Delta R^2=+0.767$, Forg.= +31.8% (final-epoch overshoot; best-epoch MSE=0.132 $\approx$ ZS). Reported for completeness; not used as primary evidence.

the $N=3$ –6 range. The ΔR^2 values at $N=4$ (+0.363) and $N=5$ (+0.767) do not show a monotone saturation pattern, but single-seed estimates are noisy and the sign is consistent.

The key dissociation is in *forgetting*: $N=3$ gives forg. -6.5 $\pm$ 4.0% (9/10 negative), $N=4$ gives +4.1% (1 seed), and $N=6$ (full unfreeze, condition B) gives -6.7 $\pm$ 6.3% (task substantially better). The $N=5$ final-epoch forgetting of +31.8% reflects a single seed that overshoot (val MSE bottoms at epoch 7 before diverging); its best-epoch MSE is 0.132, essentially unchanged from ZS—qualitatively consistent with the $N=5k$ trajectory analysis (Fig. 1, right panel). ΔR^2 and forgetting are therefore not co-determined by the same encoder layers. The top-3 of 6 layers are sufficient to restructure the encoder toward the trained objective (producing $\Delta R^2>0$), but the lower 3 layers are required for the encoder to recover pre-trained task quality. The additional ΔR^2 gain from $N=3$ to $N=6$ (+0.48 $\rightarrow$ +0.63) is modest and has much higher variance (0.115 $\rightarrow$ 0.378), consistent with the lower layers contributing primarily to task recovery rather than to the task-aligned decodability shift.

This finding sharpens the central claim: the frozen-encoder control (Appendix L) rules out data-distribution confound; the $N=3$ result rules out a whole-encoder-restructuring account.

Layer-wise drift pattern. The $N=3$ versus $N=6$ contrast reveals a within-encoder dissociation: the *top 3 layers* suffice to reorient encoder geometry toward the trained objective ($\Delta R^2>0$, 10/10 seeds, CKA $\approx$ 0.66), but the *lower 3 layers* are required for task-performance recovery ($N=3$: forg. -6.5%, 9/10 negative; $N=6$: forg. -6.7%). Both $N=3$ and $N=6$ operate in the same deeply-negative- R^2 regime; the $\Delta R^2>0$ signal at $N=3$ is a geometric reorientation, not a decodability gain—consistent with the main-text framing of frozen-probe ΔR^2 as directional evidence, not absolute evidence. (Probe-design robustness of Appendix L applies to both.)

Per-layer CKA breakdown. We computed per-layer CKA for condition B ($N=6$ full unfreeze, $n=10k$, $h=96$, 10 CUDA seeds). The bottom 3 layers are expected to drift more (lower CKA) than the top 3 under either the specialization or gradient-flow account; the data cannot distinguish these explanations without gradient-norm measurements.

Layer	CKA (mean $\pm$ std)	Group
1 (bottom)	0.441 $\pm$ 0.118	bottom 3
2	0.473 $\pm$ 0.109	bottom 3
3	0.514 $\pm$ 0.097	bottom 3
4	0.581 $\pm$ 0.088	top 3
5	0.647 $\pm$ 0.076	top 3
6 (top)	0.718 $\pm$ 0.064	top 3
Bottom 3 mean	0.476	
Top 3 mean	0.649	

Bottom layers drift more (mean CKA=0.476 vs. 0.649 for top 3), *despite* the top 3 being sufficient for $\Delta R^2>0$ ($N=3$ analysis above, 10/10 seeds). The more parsimonious explanation is standard gradient-magnitude patterns: under full backpropagation, earlier layers receive larger weight updates, producing lower CKA mechanically. Functional specialization remains an alternative hypothesis.

esis requiring a gradient-norm-per-layer comparison to distinguish from the gradient-flow null. The global CKA=0.518 at $N=6$ (Table 14) is the aggregate of these per-layer values.

O Chronos-T5-Small on M4-Monthly: Full Diagnostic Results

This appendix provides the complete results for the cross-backbone validation on Chronos-T5-Small (65M parameters, T5 seq2seq architecture with tokenised cross-entropy loss) fine-tuned on M4-Monthly (200 series, lookback=96, horizon=18). The gate-pass is 84.5% improvement over Linear baseline.

Protocol. 10 CUDA-deterministic seeds (42, 101, 123, 202, 303, 456, 777, 789, 888, 999), batch size 32, lr=1e-5, AdamW with weight decay 0.01, gradient clipping at 1.0, 20 epochs with early stopping (patience 5). Evaluation: batched median prediction (20 samples) on 200 held-out test windows.

Condition B (full fine-tune), $n=500$ (10 seeds).

	Forgetting%	CKA	ΔR^2	R^2 (PT)
Mean $\pm$ std	+1.6 $\pm$ 3.9	0.999 $\pm$ 0.001	-0.001 $\pm$ 0.002	+0.068 $\pm$ 0.025
Seeds neg/pos	4/10 neg forg	—	3/10 pos	10/10 pos

At $n=500$ the encoder is essentially unchanged (CKA $\approx$ 1). The positive R^2 (PT) confirms that Chronos’s pre-trained T5 encoder representations carry decodable forecasting structure—providing a second positive-floor anchor distinct from Moirai.

Condition B (full fine-tune), $n=10k$ (10 seeds).

	Forgetting%	CKA (ES)	ΔR^2	R^2 (PT)
Mean $\pm$ std	+0.5 $\pm$ 2.7	0.983 $\pm$ 0.021	$\approx$ 0	+0.047 $\pm$ 0.016
Seeds positive	6/10 forg	—	1/10	10/10 pos

At $n=10k$, early-stopped CKA averages 0.983 (vs. 0.999 at $n=500$). Only 4/10 seeds improve over zero-shot (best epoch $>$ 0, those 4 show CKA $\approx$ 0.95–0.97); the remaining 6/10 revert to pre-trained weights under early stopping (best_epoch=0, CKA=1.000 by definition). The strong gate-pass (84.5%) means Chronos’s encoder is already near-optimal for M4-Monthly, leaving minimal room for fine-tuning improvement. While 6/10 seeds show positive (worse) early-stopped forgetting, the magnitude (+0.5 $\pm$ 2.7%) is within the noise floor established by the frozen-encoder control (condition D: +0.6 $\pm$ 5.3%); a sign test is not statistically distinguishable from chance ($p=0.75$, two-sided binomial), confirming the 6/10 pattern reflects seed-level variation rather than systematic degradation.

Final-epoch comparison (protocol robustness). Training runs to epoch 20 regardless of early stopping; we report CKA at the final training epoch to address whether encoder stability is intrinsic or an artefact of early-stopping selection:

	Final-epoch CKA	Final-epoch Forg.%	n
$n=500$ (10 seeds)	0.968 $\pm$ 0.009	+3.1 $\pm$ 5.8	500
$n=10k$ (10 seeds)	0.955 $\pm$ 0.006	+3.4 $\pm$ 5.2	10k

All 10 seeds at $n=10k$ show final-epoch CKA $>$ 0.94 (min=0.945, max=0.964), confirming that encoder stability is *intrinsic* rather than an artefact of early-stopping selection. The mild CKA decline from 0.999 (early-stopped) to 0.955 (final-epoch) represents continued gradient pressure over 20 epochs, but remains far above the restructuring threshold observed on Moirai (CKA=0.52 on Small, 0.25 on Base). Final-epoch forgetting is +3.4 $\pm$ 5.2%—mild degradation vs. the near-zero early-stopped result, but qualitatively different from Moirai-Small’s +14% at $n=500$.

Mild CKA decline interpretation. The 0.999 $\rightarrow$ 0.955 CKA trajectory (early-stopped $\rightarrow$ final-epoch at $n=10k$) confirms a small but real gradient signal from the fine-tuning objective. We do *not* interpret this as “no drift”; rather, the drift is quantitatively minor (CKA $>$ 0.94) and functionally

benign (final-epoch forgetting +3.4% vs. Moirai-Small’s +14% at $n=500$). The value-gate predicts this: when pre-trained features are already near-optimal (84.5% advantage), the fine-tuning gradient has little room to improve and limited capacity to harm.

Condition D (frozen encoder), $n=10k$ (10 seeds). Forgetting: $+0.6 \pm 5.3\%$ (5/10 negative). CKA=1.000 (encoder frozen by definition), $\Delta R^2=0$. Decoder-only adaptation is unreliable: neither consistently better nor worse than zero-shot across seeds.

Random-init control (6 seeds). CKA=0.000 (expected: random weights have no similarity to pre-trained). $R^2(\text{PT}) < 0$ (-0.004 ± 0.005), confirming that positive-floor R^2 on the pre-trained encoder requires the learned representations, not the architecture alone. Training loss remains high (>8.0) and predictions are incoherent, confirming that Chronos’s pre-trained representations are load-bearing.

Comparison with Moirai. The contrast validates the value-gate as a regime indicator: *Moderate gate-pass* (Moirai/ETTh2, 28%): encoder restructures aggressively (CKA=0.52), probe asymmetry emerges ($\Delta R^2 = +0.67$, orthogonal < 0), forgetting is capacity- and protocol-dependent. *Strong gate-pass* (Chronos/M4, 84.5%): encoder is stable (CKA ≥ 0.95 even at final epoch), no measurable probe shift, forgetting near-zero under early stopping and mild (+3.4%) at final epoch—fine-tuning is unnecessary but not harmful. Both outcomes are correctly predicted by the gate criterion, demonstrating that the diagnostic toolkit generalises across architecturally distinct backbones. The probe asymmetry mechanism (task-specific $\Delta R^2 > 0$) requires encoder restructuring to produce a measurable signal; on Chronos, encoder stability means the asymmetry is not observable—this is a correct null prediction, not a failure of the diagnostic.

O.1 ILI Gate Check

Moirai-Small zero-shot on ILI (national illness, 966 weekly observations, 7 features, $h=24$, $\text{lookback}=104$) yields $\text{gate-pass}=57\%$ (ZS MSE=1.183, Linear MSE=2.748, per-feature normalisation by context statistics).

Full probe results (10 CUDA seeds, conditions B/D). Table 15 reports per-seed aggregate diagnostics. All 10 seeds show $\Delta R^2 > 0$ (aggregate: $+0.123 \pm 0.046$; Wilson CI [0.72, 1.00]) and forgetting < 0 ($-80.7 \pm 0.9\%$; 10/10 negative, meaning fine-tuning dramatically improves on ZS). CKA= 0.478 ± 0.022 places ILI in the restructuring regime, not the predicted intermediate regime (0.90–0.95). Condition D (frozen encoder): CKA=1.000, $\Delta R^2=0$, forgetting=0% in all 10 seeds, confirming encoder weight updates drive all changes.

Table 15: ILI per-seed results (condition B, Moirai-Small, $h=24$, $\text{lookback}=104$). Aggregate ΔR^2 is the mean across 7 univariate features.

Seed	CKA	ΔR^2	Pos. features	Forg.%	Drift
42	0.465	+0.211	6/7	−81.5	5.61
101	0.465	+0.123	3/7	−80.9	5.34
123	0.511	+0.087	3/7	−82.3	5.13
202	0.461	+0.169	5/7	−79.6	5.67
303	0.498	+0.169	7/7	−80.9	5.49
456	0.454	+0.066	5/7	−80.2	5.63
777	0.498	+0.074	4/7	−80.1	5.36
789	0.460	+0.144	6/7	−79.8	5.49
888	0.459	+0.111	4/7	−79.7	5.24
999	0.512	+0.080	4/7	−82.0	5.25
Mean±std	0.478±.022	+0.123±.046	4.7/7	−80.7±0.9	5.42±.21

Per-feature breakdown. OT (target feature) shows the largest ΔR^2 ($+0.73 \pm 0.16$, 10/10 positive); AGE 0–4 also shows consistent gain ($+0.15 \pm 0.05$, 10/10). AGE 5–24 shows slight negative signal (-0.07 ± 0.09 , 3/10 positive), consistent with the task-specificity interpretation: features less aligned with the forecasting objective show reduced or negative decodability shift.

Interpretation. The pre-registered prediction (CKA 0.90–0.95, intermediate behaviour) is partially falsified: CKA is in the restructuring regime (0.478), not intermediate. The probe-asymmetry

839 prediction ($\geq 7/10$ seeds $\Delta R^2 > 0$) is confirmed (10/10). This revises the gate-pass spectrum: the
840 restructuring-to-stability transition occurs above $\approx 80\%$ gate-pass, not at 50–70%. ILI’s low CKA
841 reflects the small dataset (966 timesteps, ≈ 450 training windows per feature) which permits aggres-
842 sive restructuring within 20 epochs.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction explicitly scope claims to the Moirai backbone family on ETT benchmarks, state that replication on additional gate-passing cells is required before claiming universality (§6), and distinguish methodology contributions from case-study findings.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Appendix A dedicates a full section to limitations: single-backbone scope, seed-count constraints, probe-head capacity (negative absolute R^2), incomplete sample-sweeps on Base/Large, and IoT scope. The conclusion also summarises key limitations.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper is empirical; it contains no theorems or formal proofs. The value-gate criterion and probe protocol are defined operationally, not as theoretical results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: §3 specifies all hyperparameters (optimizer, LR, batch size, epochs, patience), the CUDA-deterministic protocol (CUBLAS config, cudnn flags, seed list), the uni2ts patch, and hardware (A10G). Appendix sections provide per-seed results, data splits, and evaluation protocol details. The exact seed list is published.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: We cannot release code and datasets in supplementary material due to anonymisation concerns but commit to releasing them upon acceptance. ETT datasets are publicly available Zhou et al. [2021]. The paper provides sufficient detail (hyperparameters, seed lists, uni2ts patch description, evaluation protocol) to reproduce results independently.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: §3 specifies optimizer (AdamW), learning rate (10^{-4}), weight decay (10^{-2}), batch size (32), early stopping patience (3), gradient clipping (1.0), and epochs (20). Data splits (12/4/4 month chronological) are in §4.1. Full details including patch size, forecast sample count, and normalization are in Appendix D.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All multi-seed results report mean $\pm$ std with seed count stated. Sign tests with exact p -values ($p = 0.002$ for 10/10), Wilson confidence intervals (e.g., [0.72, 1.00] for 10/10), Cohen's d with 95% CI, and bootstrap p -values are reported. Error bars capture variability across CUDA-deterministic random seeds. The distinction between standard deviation and standard error is maintained throughout.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All experiments ran on a single NVIDIA A10G GPU (24GB VRAM) on an AWS g5.2xlarge instance. §3 specifies the hardware. Each $n=500$ seed takes approximately 12 minutes; each $n=10k$ seed takes approximately 35 minutes. The full 10-seed sweep at $n=10k$ requires approximately 6 hours of GPU time. Total compute for all reported experiments is under 100 A10G-hours.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research uses publicly available datasets (ETT, CICIOT2023) and pre-trained models (Moirai, Chronos, TimesFM, MOMENT) for diagnostic methodology development. No human subjects, no private data, and no dual-use concerns arise from a diagnostic methodology for representation drift.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: This paper contributes a diagnostic methodology for understanding representation drift in time-series foundation models. It does not introduce new capabilities, models, or datasets that could be misused. The methodology helps practitioners avoid deploying degraded models, which is a net positive with no foreseeable negative societal impact.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release pre-trained models, generative models, or scraped datasets. The released code implements diagnostic probes and fine-tuning scripts for existing public models and datasets, posing no misuse risk.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.

- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: All models and datasets are cited: Moirai Woo et al. [2024] (Apache-2.0), ETT Zhou et al. [2021] (MIT), Chronos Ansari et al. [2024] (Apache-2.0), TimesFM Das et al. [2024] (Apache-2.0), MOMENT Goswami et al. [2024] (MIT), CICIOT2023 (publicly released for research). The uni2ts library is cited and our patch is documented.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](#) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The supplementary code release includes the uni2ts patch, fine-tuning scripts, probe evaluation code, and reproduction instructions. The methodology is fully documented in the paper and appendix.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: This research does not involve crowdsourcing or human subjects.

1159 Guidelines:

1160 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1161 with human subjects.

1162 • Including this information in the supplemental material is fine, but if the main contri-
1163 bution of the paper involves human subjects, then as much detail as possible should
1164 be included in the main paper.

1165 • According to the NeurIPS Code of Ethics, workers involved in data collection, cura-
1166 tion, or other labor should be paid at least the minimum wage in the country of the
1167 data collector.

1168 **15. Institutional review board (IRB) approvals or equivalent for research with human**
1169 **subjects**

1170 Question: Does the paper describe potential risks incurred by study participants, whether
1171 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
1172 approvals (or an equivalent approval/review based on the requirements of your country or
1173 institution) were obtained?

1174 Answer: [N/A]

1175 Justification: This research does not involve human subjects.

1176 Guidelines:

1177 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1178 with human subjects.

1179 • Depending on the country in which research is conducted, IRB approval (or equiva-
1180 lent) may be required for any human subjects research. If you obtained IRB approval,
1181 you should clearly state this in the paper.

1182 • We recognize that the procedures for this may vary significantly between institutions
1183 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
1184 guidelines for their institution.

1185 • For initial submissions, do not include any information that would break anonymity
1186 (if applicable), such as the institution conducting the review.

1187 **16. Declaration of LLM usage**

1188 Question: Does the paper describe the usage of LLMs if it is an important, original, or
1189 non-standard component of the core methods in this research? Note that if the LLM is used
1190 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
1191 scientific rigor, or originality of the research, declaration is not required.

1192 Answer: [N/A]

1193 Justification: No LLMs are used as a component of the core methodology. The research
1194 studies time-series foundation models (Moirai, Chronos, TimesFM, MOMENT), which are
1195 not large language models.

1196 Guidelines:

1197 • The answer [N/A] means that the core method development in this research does not
1198 involve LLMs as any important, original, or non-standard components.

1199 • Please refer to our LLM policy in the NeurIPS handbook for what should or should
1200 not be described.