
Drift $\neq$ Damage: Representation Change Does Not Reliably Identify When Encoder Adaptation Helps

Anonymous Authors

Abstract

Continual- and transfer-learning practice reads the large representational change that fine-tuning a time-series foundation model (TSFM) induces as forgetting or damage. We show this reflex is misleading, and that the distinction it elides is our thesis: **drift is observed; whether encoder adaptation contributed must be intervened upon**. We ask what drift cannot answer—did *allowing* encoder adaptation contribute to the gain?—with a **frozen-encoder intervention** (B–D) behind a downstream-only **pretrained-capability inclusion criterion**. **Our headline finding is methodological: B–D must be scored on *held-out* windows**. Read on the windows model selection already uses—the natural protocol, and the one we previously reported—it systematically flatters the frozen encoder. Re-scored on a disjoint test split, three of the **19 cells** we previously reported **reverse sign** (e.g. Chronos/ETTh1 $+6.8 \rightarrow -39.2$ pp), the criterion reclassifies two others, and a third backbone added here reverses a fourth—all four the same way. Across 30 cells B–D spans -39 to $+62$ pp. Second, across four TSFMs (Moirai, Chronos, TimesFM, MOMENT), drift *magnitude* (CKA) does not reliably determine the *direction* of that effect—not even after backbone identity is absorbed as a fixed effect rather than by subsetting (App. F). Third, the intervention exposes a reproducible regime **within Moirai** that we name **pretrained-capability degradation**: in seven cells meeting the criterion (R^2_{task} 0.50–0.83) full fine-tuning leaves the model **worse than not fine-tuning at all** ($+8.7\%$ to $+46.1\%$ above zero-shot, in every seed) while a frozen encoder *improves* it (-3.5% to -16.0%); **but the criterion does not predict where**: pre-registered on seven new cells it flagged all seven and only one degraded (precision 0.14 at recall 1.00). Six of the seven are among the **least-drifted** we measure (CKA 0.67–0.89) while the cells adaptation helps most sit at 0.09–0.63; the seventh lands at 0.36, *inside* that band, so drift does not separate the two regimes at all. We establish this regime’s **existence, not its prevalence**—the third backbone contains no such cell. Two analyses from an earlier version are withdrawn and support no claim here (§4).

1 Introduction

A Moirai encoder fine-tuned on ETTh2 can have CKA with its pretrained self fall to ~ 0.3 —a $\sim 70\%$ representational shift—and yet forecast *better* on the very same task. If drift were a proxy for damage, this should not happen.

Practitioners who fine-tune released TSFM checkpoints inherit a reflex from continual learning: measure drift, read large drift as catastrophic forgetting, and reach for preservation methods (LoRA, EWC, L2-SP) [Kirkpatrick et al., 2017, Hu et al., 2022, Li et al., 2018]. This bundles three implicit claims—drift implies forgetting, forgetting implies harm, drift should be minimized—and the bundle comes apart on TSFMs. Our “ $\neq$ ” is insufficiency, not independence: drift does not *reliably imply* harm, because **downstream utility depends on the functional role of the changed representation**—not *how much* a backbone drifts but *what kind* of change it is, which one scalar cannot distinguish.

We contribute a within-model dissociation (Fig. 2), a capability check over four architecturally distinct TSFMs with controlled encoder interventions on the cells that pass it, and a diagnostic chain needing only released weights and target data (Fig. 1). On scope: the check spans four backbones, the 30 intervention cells come from Moirai, Chronos and TimesFM, and MOMENT supports no claim here. Three hypotheses, separated throughout. **H1** does drift predict utility? **H2** does drift predict how much *allowing* encoder adaptation contributes? **H3** does fine-tuning actually destroy a demonstrated pretrained capability? We give evidence against reliable monotonicity for H1 and H2; H3 yields **pretrained-capability degradation**, a reproducible regime *within Moirai*. Ordering these needs one prior result—*where* the intervention is scored—which is why that comes first below.

Our claims, strongest evidence first. (1) **B–D must be scored out-of-sample**—read on the windows model selection uses, it flatters the frozen encoder, four cells reverse sign, and the gate licensing them is vulnerable the same way. (2) **CKA drift does not reliably determine the direction of the encoder-adaptation contribution** under our evaluated protocols. (3) **A downstream-only inclusion criterion** marks where pretrained capability is demonstrated; it is not a validated detector, and we show what it would have selected prospectively. (4) **The reading is predictable without the frozen control** within Moirai, from condition B’s own held-out forgetting—a partly algebraic rule that does not transfer (§4).

2 Setup and the Diagnostic Framework

We study four TSFMs: **Moirai** [Woo et al., 2024] (encoder-only, MoE), **Chronos-T5** [Ansari et al., 2024] (encoder-decoder), **TimesFM** [Das et al., 2024] (decoder-only, $d_{\text{model}}=1280$), and **MOMENT** [Goswami et al., 2024] (encoder-only). Data is univariate ETT [Zhou et al., 2021]/M4 (lookback 96; $h=24$ cross-backbone, $h=96/192$ Moirai sweep), plus ILI (weekly influenza, 7 features, lookback 104, $h=24$) and Electricity. *Electricity is restricted to 7 of its 370 series* so its feature dimensionality—and hence the cost of a multivariate Moirai run—matches the ETT cells; this is a deliberate deviation and its numbers are not comparable to published full-Electricity results. **Each arm trains through its own objective and head.** The 21 Moirai cells—which include all seven degradation cells—and ILI train with Moirai’s *own* NLL objective through its *own* distribution head (`param_proj`) and are scored through its native sampling path (median of 20 samples); no regression head is attached, and the same holds for the 4 TimesFM cells, which train and are scored through TimesFM’s *own* point/quantile head (App. E). Only the 5 Chronos cells attach a common MSE head to encoder states, because its generative decoder admits no matched frozen-versus-adapted comparison otherwise—hence limitation (iii). What departs from the released recipes is the *data* pipeline (one dataset at a time, fixed lookback, no LOTSA packing), not the objective, head or inference path; the cost is that MSE magnitudes are *not* comparable across arms, so B–D too is read within a backbone. Two conditions run throughout: **B** (full fine-tuning) and **D** (frozen encoder, head-only). Forgetting% = $100(\text{MSE}_{\text{FT}} - \text{MSE}_{\text{ZS}})/\text{MSE}_{\text{ZS}}$; negative means improvement over zero-shot.

Three quantities must not be conflated: **representation drift** (CKA), **adaptation contribution** (B–D) and **downstream utility** (forecast MSE)—CKA measures the first, B–D the second, neither is the third. One inclusion criterion, one detector, one intervention:

Pretrained-capability inclusion criterion. $R_{\text{task}}^2 = 1 - \text{MSE}_{\text{ZS}}/\text{MSE}_{\text{Linear}}$: does the released model beat a lookback-96 linear baseline by $>20\%$? It asks whether there is capability worth preserving—not whether fine-tuning is worthwhile—and needs only released weights and the practitioner’s series. **Its role is inclusion, not prescription: it decides which cells make an intervention interpretable, and we do not claim it is a validated deployment-time decision rule**—it identifies cells with measurable zero-shot capability; we do not show that consulting it improves fine-tuning choices. A gate-fail means no demonstrated zero-shot advantage *under this protocol*, not that the pretrained representation carries no transferable information.

CKA drift. Linear CKA [Kornblith et al., 2019] between pretrained and fine-tuned backbone on mean-pooled encoder outputs over identical held-out 192-timestep windows; flagged at < 0.95 and compared only *within* a backbone, never across architectures. Crucially, **CKA measures representational change, not functional change**, which is why the chain does not stop here. We are not arguing that CKA is a poor similarity measure: the claim is against *equating* its magnitude with downstream damage, not against the measurement. Both cutoffs are *operational detectors within this protocol*, not semantic thresholds: pooling moves the CKA magnitude but not the flag (App. M, I).




Figure 1: **The chain, and why drift cannot stand in for the intervention.** (a) The three quantities the chain keeps apart. (b) All 30 cells; red marks the seven pretrained-capability degradation cells. Six are among the *least*-drifted cells we measure (CKA 0.67–0.89) while adaptation helps most at 0.09–0.63; the seventh sits at 0.36, inside that band, so no CKA cut separates the two regimes. The x -axis is not a decision boundary: CKA is comparable only *within* a backbone, and within Moirai it does not reliably order the column ($\rho=+0.181$, CI includes 0; App. F).

Table 1: **The intervention spans 101 pp; drift does not order it.** Held-out B–D (pp): negative = adaptation helps, **positive = freezing is better**. 30 cells, three backbones, 3–10 seeds; gate and B–D both on the test split. ^fgate-fails there despite passing on the selection windows; ^{*}directional ($|\text{mean}| < 2 \text{ SEM}$); [‡]B–D only (limitation iii); ^ppre-registered prospective cell. **degradation** = pretrained-capability degradation (§3); Fig. 1b plots all 30. Per-cell values App. G, statistics in App. F.

Backbone	Data (h), gate R_{task}^2	CKA	B–D held-out	Reading
Chronos [‡]	ETTh1 / ETTm2 / Weather (24), .39–.65	0.09–0.17	–39.2 / –29.5 / –21.5	adapt
Chronos [‡]	ETTh2 .46 / Electricity .14 ^f (24)	0.09–0.23	–13.0 / –2.7*	adapt
Moirai-S	ILI (24), –.03 ^f	0.455	–38.6	adapt
Moirai-B	ETTh2 (96 / 192), .29 / .55	0.40–0.46	–22.3 / –20.2	adapt
Moirai-L	ETTh2 (96), .13 ^f	0.626	–20.8	adapt
TimesFM	ETTh2 / ETTh1 (24), .70 / .47	0.25–0.36	–7.7 / –7.1	adapt
TimesFM	Weather .51 / ETTh1 .40 (24)	0.34–0.56	–1.6* / –1.4*	directional
Moirai-SP	Electricity (96 / 192), .54 / .65	0.94–0.95	–3.2 / –2.6*	adapt
Moirai-S	ETTh2, ETTm2 (96 / 192), .33–.72	0.85–0.97	+2.6 to +6.2*	directional
Moirai-BP	ETTh2, Weather (96 / 192), .48–.75	0.49–0.57	+1.3* to +12.5*	preserve
Moirai-S, LP	ETTh1 (96 / 192), .51–.72	0.36–0.84	+13.5 to +28.2	degradation
Moirai-B	ETTh1 (96 / 192), .50 / .68	0.67–0.76	+37.5 / +47.0	degradation
Moirai-S	Weather (96 / 192), .53 / .83	0.88–0.89	+53.1 / +62.1	degradation

94 **Frozen-encoder intervention (B–D).** Condition D measures how much encoder adaptation *con-*
95 *tributes* beyond head-only training; B–D is the resulting forgetting gap in percentage points, negative
96 meaning adaptation helps, and it is the only step that manipulates the encoder rather than observing
97 it. Freezing applies to encoder weights only—input/output projections stay trainable—so D is not
98 representation-frozen (its CKA runs 0.38–1.000). **B–D therefore measures whether encoder**
99 **adaptation was necessary for the gain—not that the representation change CKA measures**
100 **carried it.** Freezing also removes optimization degrees of freedom and gradient pathways, so no
101 single mechanism is isolated by this contrast.

102 3 Results: What the Intervention Separates

103 Table 1 summarizes the cross-backbone results, ordered by B–D. One structural fact first: **the**
104 **criterion is horizon-sensitive.** At $h=96$ Chronos, TimesFM and MOMENT all gate-fail (–14% to
105 –148%); re-screened at $h=24$, both backbones we re-screen flip to gate-pass ($R_{\text{task}}^2 +0.21$ to $+0.70$),
106 tracking each other dataset for dataset. Pretrained value is a property of checkpoint $\times$ dataset $\times$
107 horizon, not of an architecture; MOMENT carries no claim here (App. E).

Detection: where B–D is scored decides its sign. B–D compares two trained models and inherits their model selection: B early-stops on the evaluation windows and a frozen head’s regularizer is tuned there, so scoring B–D on those windows—the natural protocol, and the one we previously reported—is no generalization estimate. We re-scored all 19 cells on a chronologically disjoint test split, fitting and selecting no model on it (App. G; dispersions are ± 1 SEM unless a table states *std*). Three cells *reverse sign*—Chronos/ETTh1 $+6.8 \pm 2.9 \rightarrow -39.2 \pm 13.9$, and both Moirai/ETTh2 cells $(+3.8, +7.2 \rightarrow -22.3, -20.8)$ —while others merely soften. The error runs one way: read on its selection windows, B–D *flatters the frozen encoder*. TimesFM’s four cells, never re-scored, reverse a fourth the same way: four of 30, all one direction. **And so does the gate**, which also makes the seven below *retrospectively* identified. So we **pre-registered** the criterion on seven new cells—new sizes, a new dataset, no part in defining the regime—freezing each prediction from validation windows alone and committing it before any run existed. It **flagged all seven; one degraded**—precision 0.14 at recall 1.00, a base-rate predictor. It is an inclusion rule, not a detector, and we show that rather than assert it (App. G). **Intervention: pretrained-capability degradation (H3).** Our headline empirical finding is a reproducible regime *within Moirai* in which full encoder fine-tuning *destroys* the pretrained advantage the criterion demonstrated while a frozen-encoder control preserves it. Three clauses must hold, the last two in *every* seed: the cell gate-passes, B ends up worse than the un-tuned model, and D improves on it. Seven cells qualify. Moirai-Small/Weather is clearest—B lands $+37.3\%/+46.1\%$ above zero-shot error while D reaches $-15.8\%/-16.0\%$. **On ETTh1 the regime holds at all three released sizes**, Large added prospectively ($+14.1\%$ under B against -14.1% under D, 3/3 seeds)—but it is not a dataset property: Moirai-Base on Weather, run in the same batch, shows none. It tracks dataset *crossed with size*.

Robustness. Freezing the input projections too (condition H, CKA exactly 1.0) keeps the sign on all six cells it was run on, moving the magnitude by at most 5.8 pp—“freezing wins” is not the input projection re-fitting—though only **5/6** beat zero-shot under the stricter H (App. H). Chronos/M4 (gate 84.5%) anchors the stable end (App. L). **A third backbone.** TimesFM-2.5 under its *own* recipe—no head attached, so both forgetting terms are clean—gate-passes on four datasets, and **none is a degradation cell**: B improves on zero-shot in all four, B–D runs -7.7 to -1.4 . The regime is Moirai’s, at horizons where TimesFM gate-fails: scope, not generality. It also gives the measurement claim its sharpest form: condition D’s transformer stack is **bit-identical** after training (ℓ_2 drift exactly 0.00) and still reads CKA 0.38–0.89—all of it the input tokenizer—so low CKA cannot by itself mean a damaged encoder (App. E).

Scope: drift does not flag the degradation cells. Six of the seven sit at CKA 0.67–0.89 while the cells adaptation helps most sit at 0.17 and 0.46—the reflex would have cleared the damaging cells and flagged the safe ones—and the seventh sits at 0.36, among them: no CKA cut separates the regimes. Nor does it order the held-out column *within* a backbone— $\rho = +0.181$ over 21 Moirai cells, and the slope stays unpinned when backbone identity is absorbed as a fixed effect instead ($n=30$, App. F). Within one model they come apart directly (Fig. 2, App. A).

4 Discussion

What this is and is not. The sign reversals are a generalization result—no model is fitted or selected on test. The degradation regime is *identified* using test data, so it establishes existence, not prevalence; the pre-registered arm is what bounds its detection accuracy (App. G). Runs were not blinded and two pre-specified predictions were falsified.

A cheaper proxy, where it holds. Since Moirai’s frozen arm barely moves, condition B’s own forgetting agrees in sign with B–D in 15 of 21 cells ($\rho = +0.897$)—one run, no control, but partly algebraic and it does not transfer (App. J).

What follows for practice. Check the checkpoint beats a trivial baseline before worrying about preserving it; do not read drift as damage; and score frozen-versus-adapted on *held-out* windows—on selection windows it favours freezing.

Limitations (related work: App. K). (i) Two analyses from an earlier version are **withdrawn**; B–D shares no machinery with either (App. B, I). (ii) Cells are unequally powered (3–10 seeds) and **not independent**—datasets, horizons and sizes recur—so cross-cell statistics are descriptive. (iii) Chronos carries no per-condition forgetting ($\ddagger$, Table 1), and the seven degradation cells are Moirai-only on two datasets. (iv) Condition D leaves input projections trainable (App. H).

References

- Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, Jasper Zschiegner, Danielle C. Maddix, Hao Wang, Michael W. Mahoney, Kari Torkkola, Andrew Gordon Wilson, Michael Bohlke-Schneider, and Yuyang Wang. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024.
- Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *International Conference on Machine Learning (ICML)*, 2024.
- Mononito Goswami, Konrad Szafer, Arjun Choudhry, Yifu Cai, Shuo Li, and Artur Dubrawski. MOMENT: A family of open time-series foundation models. In *International Conference on Machine Learning (ICML)*, 2024.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning (ICML)*, pages 3519–3529, 2019.
- Ananya Kumar, Aditi Raghunathan, Robbie Jones, Tengyu Ma, and Percy Liang. Fine-tuning can distort pretrained features and underperform out-of-distribution. In *International Conference on Learning Representations (ICLR)*, 2022.
- Xuhong Li, Yves Grandvalet, and Franck Davoine. Explicit inductive bias for transfer learning with convolutional networks. In *International Conference on Machine Learning (ICML)*, 2018.
- Zhizhong Li and Derek Hoiem. Learning without forgetting. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(12):2935–2947, 2018.
- Yun Luo, Zhen Yang, Fandong Meng, Yafu Li, Jie Zhou, and Yue Zhang. An empirical study of catastrophic forgetting in large language models during continual fine-tuning. *arXiv preprint arXiv:2308.08747*, 2023.
- Behnam Neyshabur, Hanie Sedghi, and Chiyuan Zhang. What is being transferred in transfer learning? In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- Michał Wiliński, Mononito Goswami, Willa Potosnak, Nina Żukowska, and Artur Dubrawski. Exploring representations and interventions in time series foundation models. In *International Conference on Machine Learning (ICML)*, 2025.
- Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *International Conference on Machine Learning (ICML)*, 2024.
- Haoyi Zhou, Shanghang Zhang, Jieqi Peng, Shuai Zhang, Jianxin Li, Hui Xiong, and Wancai Zhang. Informer: Beyond efficient transformer for long sequence time-series forecasting. In *AAAI Conference on Artificial Intelligence*, 2021.

204 A Within-Model Dissociation Sweep

205 Every point shares a backbone, a dataset, a horizon and a protocol; only the fine-tuning sample size
 206 varies. Increasing n moves drift monotonically while the task outcome moves non-monotonically:
 207 CKA falls steadily from 0.949 at $n=500$ to 0.518 at $n=10k$, while forgetting goes +14.0, +13.4,
 208 -2.9, +6.6, -5.3%. The best cell in the sweep is the *most-drifted* one. The $n=5k$ variance (+6.6%,
 209 std 10.0, 7 seeds) is bimodal—4 seeds near -0.1%, 3 at +15.5% with an epoch-1 overshoot—so we
 210 treat that point as unresolved rather than as evidence for either direction.

211 This is a *dissociation, not an intervention on drift*. Sample size is an uncontrolled causal variable:
 212 raising it changes supervision, optimization trajectory and overfitting, not drift alone. The sweep
 213 therefore shows that drift and outcome need not move together within one model; it cannot show that
 214 drift itself is causally inert. The frozen-encoder experiment is the intervention, and it is what the
 215 paper’s claims rest on.




Figure 2: **Within-model dissociation.** Moirai-Small/ETTh2 ($h=96$, cond. B) across n : as n rises CKA drift is monotonic (top) while forgetting is non-monotonic and sign-flipping (bottom), improving at maximum drift: **drift and downstream outcome can move independently**. Sample size is uncontrolled here, so this demonstrates **non-monotonic co-movement, not causal independence**: it is a dissociation, not an intervention on drift. Error bars ± 1 SD; $n=10k$ is 10-seed CUDA-deterministic.

216 B The Probe Evidence We Retract

217 **This appendix records a retracted result.** It reports the probe statistics as we originally computed
 218 them, unchanged, so that they can be audited against the re-analysis in Appendix I that withdraws
 219 them. Nothing in the main results depends on this section; we keep it rather than deleting the numbers.

220 Table 2 gives every probe used in the asymmetry test on Moirai-Small/ETTh2, with dispersion and
 221 seed counts. We classified each ΔR^2 by an effect-size heuristic—whether the mean exceeds one
 222 seed-level SD—not by a significance test. Four auxiliary probes decrease; one is indistinguishable
 223 from zero; none increase. The signature as stated—trained increases **and** no auxiliary increases—is
 224 satisfied without relying on the forecast-48 result in either direction. What it is *not* robust to is the
 225 operating point, which is the subject of Appendix I.

226 **Why these five.** The auxiliary set is chosen to vary along two axes a single probe family would
 227 confound: *what* is decoded (low-order statistics vs. forecasts) and *how* it is decoded (linear Ridge
 228 vs. gradient-boosted). The delta1 probe matters specifically because its Ridge version is misleading:
 229 the zero-shot Ridge baseline on next-step deltas is deeply negative ($R^2 \approx -2.3$), so any encoder
 230 shift inflates ΔR^2 (Ridge gives $+1.04 \pm 0.22$). A GBM probe establishes a genuine positive floor

Table 2: Probe ΔR^2 on Moirai-Small/ETTh2, $h=96$, $n=10k$. Auxiliary targets span statistical properties (mean, variance, lag-1 autocorrelation), an alternative forecast horizon (forecast-48), and a positive-floor readout (GBM delta1).

Probe target	Trained?	Seeds	ΔR^2	Classification
Forecast-96 (objective)	yes	10	$+0.67 \pm 0.23$ (10/10 pos)	increase
Mean	no	10	-0.033 ± 0.019 (10/10 neg)	decrease
Variance	no	10	-0.048 ± 0.030 (9/10 neg)	decrease
Lag-1 autocorrelation	no	10	-0.052 ± 0.041 (9/10 neg)	decrease
Delta1 (GBM, positive floor)	no	4	-0.056 ± 0.030	decrease
Forecast-48 (sub-horizon)	no	4	-0.001 ± 0.163	indist. from zero

($R^2(\text{PT})=+0.059$) and on that floor decodability *falls*. We report the GBM result for this reason, and note that this is the only positive-floor untrained readout available to us—a reviewer-visible weakness rather than a strength. Five probes cannot establish that no specialization occurs along an unmeasured axis (calibration, periodicity, tail behaviour); the asymmetry test is evidence, not proof.

C ILI Cell Detail

The gate value we previously reported for this cell does not reproduce. Earlier versions quoted gate-pass =57% (ZS MSE 1.183, Linear MSE 2.748) from a run set that was not retained. Re-computed from the runs this paper actually reports (10 seeds, results/v40_ili_heldout) with the same definition and the same per-window normalisation, the gate is +0.139 on validation (ZS 1.893, linear 2.206) and **-0.033** on the test segment (ZS 2.479, linear 2.664): **ILI gate-fails on held-out windows**, and we withdraw the 57%. The per-feature picture is why the mean lands where it does—five of seven features are positive but only two clear 0.20, and NUM. OF PROVIDERS is -0.86 because that series is close to a straight line on the test segment, which linear extrapolation fits almost exactly. ILI is not a degradation cell and no claim in §3 rests on its gate; what changes is that it can no longer be cited as a moderate-gate cell, and the “falsified prediction” below is correspondingly weaker evidence—a heavily restructuring encoder on a cell with *no* demonstrated zero-shot advantage is much less surprising than one on a moderate-gate cell.

Status of this cell. On ILI the encoder restructures heavily (CKA 0.455 ± 0.008) and full fine-tuning improves by $-80.0 \pm 0.6\%$ (10/10), falsifying a prediction specified beforehand (0.90–0.95 CKA expected). We previously reported B–D= -80.7 pp here, i.e. that encoder adaptation supplied *all* of the gain. That is withdrawn. The frozen control it was measured against had never been trained: the script assigned the zero-shot MSE as the fine-tuned MSE and wrote a forgetting of exactly 0.0, so the “0.0% in all 10 seeds” was one constant, not ten measurements. Re-run so that condition D actually trains — encoder frozen, the 3.2M parameters outside it updated — the gap is -62.6 ± 1.4 on validation and -38.6 ± 1.3 held-out (10/10 negative on both). D reaches CKA 0.76–0.90, not 1.000, because the projections feeding the encoder remain trainable. Condition B is unaffected and reproduces the published -80.7 ± 0.9 and CKA 0.478 to within seed noise, so the entire discrepancy was the untrained control. Note also that both references are weak in absolute terms (ZS MSE 1.893, linear 2.206 on validation).

The cell is Moirai-Small on national influenza-like illness: 966 weekly observations, 7 features, $h=24$, lookback 104, split 579/193/194 chronologically. Condition D, re-run so that it trains, gives CKA 0.827 ± 0.013 and forgetting $-17.4 \pm 1.3\%$ on validation ($-12.2 \pm 1.3\%$ held-out), 10/10 negative: head-only training recovers a substantial share of the improvement, so the B–D gap is -62.6 pp on validation and -38.6 pp held-out, not the -80.7 previously reported.

Per-feature asymmetry. The forecast target OT shows the largest gain ($+0.73 \pm 0.16$, 10/10 positive) and AGE 0–4 also gains consistently ($+0.15 \pm 0.05$, 10/10). AGE 5–24 is negative (-0.07 ± 0.09 , 3/10 positive). Decodability thus improves on the target and degrades on a series the head was not optimizing—the same signature as the ETTh2 trained-vs-auxiliary split, obtained here within a single cell.

Falsified prediction. We specified CKA 0.90–0.95 for ILI in advance, on the expectation that a moderate gate implies moderate drift. The observed 0.456 falsifies this: gate strength does not linearly

Table 3: ILI per-seed results (Moirai-Small, $h=24$, lookback 104), both conditions, regenerated for this revision. Each value is the mean across the 7 univariate features. The ΔR^2 column of the earlier version of this table is dropped with the retracted probe axis; condition D is shown because it previously did not exist as a measurement.

Seed	B (full fine-tuning)		D (frozen encoder)	
	CKA	Forg.% (val / held-out)	CKA	Forg.% (val / held-out)
42	0.458	−78.6 / −51.5	0.804	−16.2 / −14.6
101	0.465	−81.4 / −50.1	0.898	−13.5 / −6.7
123	0.476	−80.0 / −48.9	0.756	−27.6 / −19.5
202	0.457	−81.4 / −52.1	0.791	−19.7 / −13.6
303	0.409	−78.1 / −49.0	0.823	−18.3 / −11.5
456	0.463	−75.5 / −50.0	0.851	−15.4 / −13.2
777	0.437	−81.6 / −52.5	0.822	−18.0 / −12.4
789	0.449	−80.3 / −49.5	0.865	−12.7 / −4.9
888	0.505	−82.0 / −50.8	0.848	−16.1 / −12.8
999	0.437	−80.8 / −53.4	0.807	−16.1 / −12.7
Mean	0.456	−80.0 / −50.8	0.827	−17.4 / −12.2

predict drift magnitude, and the restructuring-to-stability transition sits nearer $\approx 80\%$ gate than the 50–70% we assumed. We report this because it is the kind of prediction the chain is supposed to expose.

D Full Preservation Spectrum

What this appendix is for. It is a *secondary*, observational result: it shows that raising representational similarity does not raise utility, which is evidence against the reflex’s prescriptive half. It is **not** evidence for this paper’s causal claim—that rests entirely on the frozen-encoder intervention B–D in §3, because these methods preserve representations by different mechanisms and so are not a controlled manipulation of one variable. The main text quotes $h=96$ only. Table 4 gives both horizons and all methods, including the NLL+SupCon variant. These methods preserve representations by *different mechanisms*—penalty (L2-SP, EWC), reparameterization (LoRA), and complete freezing—so CKA is a common observational axis, not a common intervention, and at $n=4$ –7 the correlations here are descriptive only. We stress that **LoRA does well**: the paper’s claim is that drift magnitude fails to *order* these methods, not that preservation is a bad idea or that fine-tuning should be avoided. **A correction to an earlier version.** The figures we previously reported for this spectrum were not matched to their baseline. The mitigation runs use $n=1000$ while the full-fine-tuning reference we quoted against them was an $n=500$, 10-seed run, and the mitigation forgetting percentages divided by a single shared zero-shot (0.1259) rather than each run’s own, which sits between 0.1274 and 0.1310; that inflated every one of them by 1.2–3.4 pp. Corrected—all at $n=1000$, $h=96$, each run against its own zero-shot—full fine-tuning is CKA .936 at +13.4%, against L2-SP $\lambda=0.01$ (.939, +16.0), L2-SP $\lambda=0.1$ (.965, +13.4), EWC $\lambda=1000$ (.972, +14.9) and LoRA (.979, −9.2).

The claim this displaces is the strong one: it is *not* true that every weight-anchoring configuration forecasts worse. L2-SP $\lambda=0.1$ raises CKA substantially (.936 $\rightarrow$.965) and lands on full fine-tuning to within 0.02 pp. What the corrected spectrum shows is the weaker and more useful thing: moving up the CKA axis buys nothing predictable—one anchoring setting is worse, one is a dead heat, a third is worse again, and the reparameterizing method is far better while being the most CKA-preserving of the four. Representation preservation does not predict utility here. At $h=192$ no $n=1000$ full-fine-tuning baseline exists in our runs, so we report that horizon’s methods without a matched reference rather than against the $n=500$ one.

Why we do not report a single CKA–utility correlation. A natural request is a correlation over Table 4 rather than a verbal reading of seven points. We computed it, and the answer is instructive: *the sign depends on which mechanisms are pooled*.

Table 4: Mitigation spectrum on ETTh2 forecasting. Methods ordered by representation preservation (CKA, $h=96$). LoRA preserves representations (CKA ≈ 0.98) while *improving* over zero-shot. Mean $\pm$ std; Forgetting% = (fine-tuned MSE – zero-shot MSE) / zero-shot MSE $\times$ 100. **Two caveats on comparability, added in revision.** Forgetting% divides by the single shared zero-shot printed in each horizon header, not by each run’s own zero-shot; the anchoring and LoRA runs measure their own at 0.127–0.131, so their figures here are inflated by 1.2–3.4 pp. And the rows are not matched on training-set size: B and C use $n=500$ training samples while L2-SP, EWC and LoRA use $n=1000$ (the n column counts seeds, not samples). Matched at $n=1000$, $h=96$, each against its own zero-shot: full fine-tuning .936/+13.4, L2-SP $\lambda=0.01$.939/+16.0, L2-SP $\lambda=0.1$.965/+13.4, EWC .972/+14.9, LoRA .979/–9.2. On that comparison L2-SP $\lambda=0.1$ ties full fine-tuning rather than losing to it, so anchoring is not uniformly harmful; the dissociation the table is cited for—CKA not predicting utility—is unaffected.

Method	$h = 96$ (ZS MSE: 0.126)				$h = 192$ (ZS MSE: 0.158)				n
	MSE	Forg.%	CKA	Drift	MSE	Forg.%	CKA	Drift	
B: Full fine-tune	.140 $\pm$.014	+11.1	.937 $\pm$.028	3.96 $\pm$.56	.183 $\pm$.020	+15.7	.970 $\pm$.009	4.12 $\pm$.32	10
F: L2-SP ($\lambda=0.01$)	.150 $\pm$.006	+19.0	.939 $\pm$.023	4.98 $\pm$.34	.179 $\pm$.009	+13.3	.945 $\pm$.007	4.84 $\pm$.22	5
C: NLL+SupCon	.138 $\pm$.010	+10.1	.949 $\pm$.014	3.72 $\pm$.43	.178 $\pm$.011	+12.9	.964 $\pm$.015	3.97 $\pm$.34	10
F: L2-SP ($\lambda=0.1$)	.147 $\pm$.003	+16.7	.965 $\pm$.006	2.94 $\pm$.08	.181 $\pm$.003	+14.6	.980 $\pm$.002	2.84 $\pm$.02	5
G: EWC ($\lambda=1000$)	.149 $\pm$.015	+18.3	.972 $\pm$.004	3.29 $\pm$.79	.188 $\pm$.004	+19.0	.978 $\pm$.006	4.46 $\pm$.38	5
E: LoRA ($r=8$) [†]	.116 $\pm$.001	–7.9	.979 $\pm$.006	0.00 $\pm$.00	.157 $\pm$.002	–0.6	.986 $\pm$.005	0.00 $\pm$.00	5
D: Frozen encoder	.126 $\pm$.006	+0.5	1.000 $\pm$.000	1.11 $\pm$.15	.155 $\pm$.002	–1.7	1.000 $\pm$.000	1.30 $\pm$.18	10

[†] LoRA ($\alpha=16$) adds low-rank updates that do not modify base encoder parameters; ℓ_2 drift is exactly zero by construction.

Conditions included	$h=96$		$h=192$	
	Pearson r	Spearman ρ	Pearson r	Spearman ρ
Weight-anchoring only ($n=4$)	+0.46	+0.40	+0.64	+0.40
All seven conditions ($n=7$)	–0.58	–0.54	–0.60	–0.43

Forgetting is negative-is-better, so a *positive* coefficient means higher CKA accompanies worse utility. Among weight-anchoring methods the sign is positive at both horizons; pooling all seven flips it negative, because LoRA and the frozen encoder sit at the highest CKA and perform well. Neither coefficient is meaningful on its own at $n=4$ or $n=7$, and we claim no significance. The point is the flip: CKA is a common *observational* axis across these methods, not a common intervention, so no single correlation summarizes the table. This is the same reason we present the spectrum as a mechanism-separated comparison rather than a regression.

E Chronos and TimesFM Intervention Detail

Chronos-T5-Small, moderate gate ($n=8k$, conditions B and D). This pair is the main text’s first inversion. An earlier version of this table reported it from runs that gave both conditions the same 30-epoch budget. That budget is generous for B, which early-stops with its best epoch at 1–2, but it truncates D badly: every D seed restored “best epoch 30 of 30” with validation loss still falling nearly linearly, so the apparent gap measured the budget rather than the intervention. We re-ran both cells with each condition trained to its own convergence.

Because in condition D the encoder is frozen *and* in `eval()` mode, its mean-pooled output is a deterministic function of the context window, and the head is a single linear map on that vector. Caching the pooled features is therefore exact, not an approximation, and it makes convergence affordable. We then estimate the best head two independent ways: a closed-form ridge optimum over 16 log-spaced penalties (the headline, since a closed form admits no epoch budget at all; the penalty is selected on the same eval windows B early-stops on, so both sides receive eval-set selection, and the optimum was interior to the grid in all six runs), and a protocol-matched AdamW run to early stopping. The two agree to about 1 pp everywhere, which is the evidence that the number is a property of the frozen features rather than of the optimizer.

One further consequence of the re-run is why Table 1 reports B–D for these cells but no per-condition forgetting. Forgetting is a ratio against the zero-shot reference, and for Chronos the fine-tuned

Table 5: Chronos-T5-Small, $h=24$, $n=8k$, matched seeds and eval windows. MSE is per-element validation MSE; $B-D=(MSE_B - MSE_D)/MSE_{ZS} \times 100$, so the zero-shot reference cancels and only the two same-head models are compared. All six D runs converged (early stopping fired) and all six ridge optima were interior.

Cell	Seed	MSE _B	D (ridge optimum)		B-D (pp)		
			MSE _D	α	ridge	AdamW	CKA (B)
ETTh1	42	1.598	1.580	0.3	+1.2	+0.5	0.179
	43	1.471	1.364	0.03	+7.5	+5.9	0.158
	44	1.509	1.365	0.1	+11.4	+11.1	0.171
mean $\pm$ std					+6.7 $\pm$ 5.1	+5.8 $\pm$ 5.3	0.169
ETTh2	42	0.998	1.133	0.03	-13.3	-14.8	0.094
	43	0.842	1.053	0.03	-23.4	-25.0	0.102
	44	0.882	1.078	0.03	-22.8	-21.7	0.087
mean $\pm$ std					-19.9 $\pm$ 5.7	-20.5 $\pm$ 5.2	0.094

model is scored through the trained MSE head while that reference is scored through Chronos’s generative decoder—so the ratio charges the output-head swap to forgetting and is not a measurement of adaptation at all. B-D is immune because both conditions share the head and the reference cancels algebraically. An earlier version of the table did quote those per-condition numbers, and the direction of the effect they implied does not survive: measured condition B is +12.6% on ETTh1 and -1.6% on ETTh2. It also quoted a Chronos CKA range of 0.886–0.949, which we have traced to the Moirai-Small sample-size sweep; the measured Chronos values are 0.087–0.179.

On the validation windows every ETTh1 seed sits above every ETTh2 seed, which an earlier version read as an opposite-signed pair. That separation does not survive held-out scoring: on the test windows both cells are negative (ETTh1 -39.1 ± 24.1 , ETTh2 -12.7 ± 1.0 , 3/3 each), so we no longer report a sign contrast between them (App. G). What the two cells still show is that CKA does not order the outcome—ETTh1 drifts *less* (0.158–0.179 against 0.087–0.102) and loses *more* to freezing.

Chronos-T5-Small, strong gate. On M4-Monthly (gate 84.5%, 10 seeds) condition B gives +0.5% and condition D +0.6% (10 seeds; std 2.7 and 5.3); the B-D gap is within the frozen-control noise floor, and a sign test on the 6/10 positive-forgetting seeds is not distinguishable from chance ($p=0.75$, two-sided binomial).

Cross-dataset transfer. Fine-tuning Moirai-Small on ETTh2 improves ETTh1 ($-7.7 \pm 2.5\%$) and ETTm2 ($-3.3 \pm 1.0\%$), 3 seeds, condition B; the frozen control shows near-zero out-of-distribution effect.

Pretraining overlap. Moirai’s LOTSA corpus excludes ETT; M4 is in Chronos’s pretraining data (which is why the M4 cell functions as an overlap control); overlap for TimesFM and MOMENT could not be established from public documentation.

TimesFM-2.5-200M at $h=96$. Five seeds on ETTh1 across conditions B and D give stability throughout (CKA 0.995–0.998, $\Delta R^2 \approx 0$, no asymmetry), and the checkpoint gate-fails at this horizon ($R^2_{\text{task}} -0.139$ on ETTh1, -0.138 on ETTh2), so our own criterion does not license an intervention there. Per-seed forgetting figures for the two conditions were not retained, so this pair is not one of the 19 re-scored cells of Table 1 rather than being entered with an estimate.

TimesFM-2.5-200M at $h=24$: the criterion. Because the criterion is horizon-sensitive, we re-screened TimesFM at $h=24$ before running any intervention, on the *same* held-out window set the Chronos arm uses—univariate OT, lookback 96, the 200-window seed-0 test subsample—so the criterion, the zero-shot reference and B-D all read one window set. All five datasets pass, and they track Chronos dataset for dataset, which is evidence that R^2_{task} is measuring a property of the dataset $\times$ horizon rather than of the backbone:

R^2_{task} at $h=24$ (test)	ETTm2	Weather	ETTh2	ETTh1	Electricity
TimesFM-2.5-200M	+0.703	+0.506	+0.472	+0.401	+0.214
Chronos-T5-Small	+0.651	+0.494	+0.456	+0.389	+0.138

Both arms are scored by `scripts/gate_all_cells.py`, one code path for all four arms, cached to `results/gate_test_side.json`.

TimesFM-2.5-200M at $h=24$: the intervention. Having licensed it, we run the same conditions as the Moirai arm (`scripts/finetune_timesfm.py`): **B** trains everything; **D** freezes the 20-layer stack (196.71M of 231.29M parameters) and leaves the input tokenizer and output projections trainable, exactly as Moirai’s **D** leaves `in_proj` trainable. Hyperparameters are the ones the other arms use— $\text{lr } 10^{-4}$, AdamW, grad-norm clip 1.0, 20 epochs, patience 7, 1000 training windows—and are identical for **B** and **D**, since **B**–**D** would be meaningless otherwise. After restoring the best epoch we *assert* the frozen modules are bit-identical to the released checkpoint; the run aborts rather than mislabel a condition.

Two design points matter for reading these cells. First, no head is attached: the loss is TimesFM’s own output structure (MSE on the mean channel plus pinball loss on the nine decile channels the released output `projection_point` emits), and zero-shot, **B** and **D** are all scored through the same native path. **Per-condition forgetting is therefore a clean measurement here**, unlike Chronos (limitation iii), so this arm re-tests the `forB` rule on a second backbone. Second, the released torch API cannot be trained through—`module.decode()` is wrapped in `torch.no_grad()` and returns detached NumPy—so training uses a differentiable re-implementation of the prefill path. That path is legitimate only because at $h=24$ the autoregressive loop runs $(\text{max_horizon} - 1)/o = 0$ times with $o=128$: one forward pass is the whole forecast. `verify_native_path()` asserts agreement with `model.forecast()` on real windows before any training (max relative deviation 6.5×10^{-7} on MPS, 0.0 on CPU), so a re-implementation error cannot reach the paper silently. The continuous quantile head (27.85M) rewrites only non-median channels at inference and receives no gradient, so condition **B** updates 203.44M of the 231.29M parameters.

What the four cells show. We ran the four highest-gate datasets, 3 seeds $\times$ {**B**, **D**} each, in an order fixed before any run by correspondence to the degradation cells (ETTh1 first, then Weather, then ETTm2, ETTh2). **All four land in the *adapt* half of Table 1, and none is a degradation cell:** condition **B** improves on zero-shot everywhere (`forB` -51.8% to -3.0%), which is the first clause of degradation failing, so the regime does not extend to this backbone at $h=24$. That is a scope limit on H3 (limitation iv), not a refutation of it: the degradation cells sit at $h=96/192$, where TimesFM *gate-fails* and our own criterion forbids the comparison.

TimesFM, $h=24$	gate	CKA (B)	CKA (D)	<code>for_B</code>	<code>for_D</code>	B – D val	B – D held-out
ETTm2	+0.703	0.19–0.52	0.48–0.55	-51.8 ± 2.8	-44.1 ± 1.6	-12.2 ± 12.4	-7.7 ± 2.6
Weather	+0.506	0.24–0.40	0.62–0.71	-7.8 ± 1.6	-6.1 ± 0.5	-4.8 ± 0.6	-1.6 ± 1.5
ETTh2	+0.472	0.22–0.28	0.38–0.50	-20.4 ± 1.4	-13.3 ± 1.9	$+0.2 \pm 2.9$	-7.1 ± 3.1
ETTh1	+0.401	0.44–0.64	0.80–0.89	-3.0 ± 2.5	-1.6 ± 0.9	-5.4 ± 0.9	-1.4 ± 2.8

Three things the arm does establish, none of which was available from Moirai or Chronos.

(a) The held-out reversal replicates on fresh cells. These four cells were scored on both splits from the same runs—nothing was re-scored after the fact—and TimesFM/ETTh2 *reverses sign*, $+0.2$ on validation against -7.1 held out. It reverses in the *same* direction as the three Moirai/Chronos reversals: validation favoured the frozen encoder, held-out data does not. Four of our 30 cells reverse and all four reverse that way, which is the paper’s central methodological claim (§3) reproduced on a backbone it was not derived from.

(b) A stack that is provably unchanged can read CKA 0.38. Condition **D**’s ℓ_2 weight drift over `stacked_xf` is exactly 0.00—the bit-identity assertion above—yet its measured CKA runs 0.38–0.89. Every point of that movement comes from the input tokenizer changing what the frozen stack is shown. This is the sharpest form of the paper’s measurement claim: CKA registers representational change on a network whose weights did not move at all, so a low CKA cannot by itself be read as the encoder having been damaged. It also bounds limitation (v) from the other side—**D** is not representation-frozen anywhere, and here we can say exactly why.

(c) **Drift does not order the four.** CKA (B) spans 0.19–0.64, comparable to the degradation cells’ drift, while held-out B–D stays within 8 pp of zero; the ranking is not monotone ($\rho = +0.400$, $n=4$, CI $[-0.778, +1.000]$), and the *most*-drifted cell (ETTh2, CKA 0.22–0.28) is among the two where adaptation helps most. Because no head is attached, forg_B is a clean measurement here, so this arm also re-tests the §4 rule: its *sign* transfers— $\text{forg}_B < 0$ and $B-D < 0$ in 4 of 4 cells—but its stated mechanism does not, since forg_D varies here almost as much as forg_B (SD 19.2 against 22.0, versus 8.0 against 26.8 on Moirai). The rule is an empirical property of a nearly-constant frozen arm, and that property is Moirai’s, not a general one.

We do not re-screen MOMENT: its $h=96$ margin is -147.8% , and its released checkpoint has no pretrained forecasting head, so the zero-shot term the criterion needs does not exist without fitting a probe—which would make the “zero-shot advantage” an artifact of the probe. MOMENT therefore supports no claim in this paper.

F Cross-Cell Statistics, and a Pooled Analysis We Do Not License

Across the 30 cells of Table 1 we ask whether either observational diagnostic orders the held-out intervention. Spearman ρ with a 10,000-sample bootstrap CI over cells:

Grouping	n	CKA range	ρ	95% CI
Pooled across backbones	30	0.09–0.97	+0.576	[+0.281, +0.759]
Moirai only	21	0.36–0.97	+0.181	[−0.321, +0.610]
Chronos only	5	0.09–0.23	+0.100	[−1.000, +1.000]
TimesFM only	4	0.25–0.56	+0.400	[−0.778, +1.000]
ℓ_2 weight drift (pooled)	24	—	+0.075	[−0.348, +0.512]

The ℓ_2 row covers 24 cells rather than 30 because the Chronos runs did not retain a weight-drift measurement. **A correction to an earlier version:** that row previously read $n=19$, $\rho = +0.044$, which was in fact the Moirai-only figure ($n=13$, CI $[-0.545, +0.651]$) carrying a “pooled” label. The pooled value is -0.301 ; both include zero, so the reading does not change, but the label was wrong and the sign was not ours to report.

Absorbing the backbone instead of subsetting it. The within-backbone rule can be read as convenient—we adopt the split that removes the significant relationship. So we also fit the pooled cells directly, with backbone identity absorbed as a fixed effect rather than by subsetting, and CKA entering as a continuous regressor (`scripts/cka_fixed_effects.py`):

Model for held-out B–D	n	b_1 (pp per unit CKA)	95% CI
$b_0 + b_1$ CKA (no backbone term)	30	+47.6	—
$b_0 + b_1$ CKA + backbone FE	30	+39.5	[−14.7, +98.3]
+ CKA $\times$ backbone	30	+40.4	[−18.6, +106.7]

CI’s are a 20,000-draw bootstrap resampling *clusters*, a cluster being a (backbone, dataset) pair, because datasets and checkpoints recur across horizons and sizes and an unclustered interval would be far too narrow; draws that leave a backbone unrepresented are discarded and the retained count is printed by the script. **Both adjusted intervals include zero.** The pooled rank correlation therefore does not survive absorbing backbone identity, and it is not the within-backbone restriction that removes it—the same conclusion follows from a model fitted on all 30 cells at once. What the data cannot do is pin the within-backbone slope down; that is weaker than, and must not be read as, a claim that CKA carries no information.

The pooled row is not a licensed analysis, and it points the other way. Pooled, CKA orders the intervention with a CI excluding zero ($p=0.0001$). We do not read it as evidence, for the reason stated in §2: CKA magnitude is representation-dependent, so values from different architectures are not on a common scale, and we compare them only within a backbone. That restriction was fixed before these cells were run, not chosen after seeing this table. The pooled figure is largely what an architecture confound would produce here: every Chronos cell sits at CKA 0.09–0.23 and every Moirai cell at 0.40–0.97, so pooling those two recovers the backbone split rather than a drift–utility

relationship. The third backbone is the honest complication: TimesFM’s four cells span 0.25–0.56 and therefore *straddle* the Chronos and Moirai ranges instead of forming a third stratum, so the pooled correlation is not purely a between-backbone artifact. It is also the group where drift is largest relative to outcome—four cells at CKA 0.25–0.56 whose B–D all sit within 8 pp of zero—so adding it pushes the pooled ρ up while adding cells that individually contradict the reflex. Neither observation licenses reading the pooled row as evidence; both are why we print it.

We report it anyway because a reader can compute it in a minute, and because the honest position is that our central claim survives on a protocol restriction rather than because no correlation can be found in the numbers. Within Moirai, the only group with enough cells and enough CKA range to test, $\rho = +0.181$ with a CI that includes zero—but this is a small sample, and we do not claim CKA is uninformative, only that it does not reliably order the intervention at this resolution. Within TimesFM, the one other backbone with a clean forgetting measurement, ρ is $+0.400$ on 4 cells with a CI spanning almost the whole range—too few cells to test anything, and reported so that the arm’s failure to help is on the record rather than omitted. The ℓ_2 drift result is cleaner: it orders nothing even pooled ($\rho = -0.301$, CI $[-0.731, +0.249]$), which is the more robust form of the claim that observational summaries of fine-tuning do not substitute for running the intervention.

G Held-Out Re-Scoring of the Intervention

Why this appendix exists. Every B–D we previously reported was read on the evaluation windows that model selection already used: condition B early-stops on them, and for the converged Chronos runs the frozen head’s ridge penalty is chosen on them. That makes B–D a controlled within-cell contrast but not an estimate of generalization. Each training script in this work already carves a chronologically disjoint test split and then discards it; here we score it.

Two kinds of evidence, and only one of them is prospective. These are separate claims and the paper does not conflate them. (A) *The protocol finding is genuinely out of sample.* No model is fitted or selected on the test split, so the sign reversals in §3 are a generalization result. (B) *The degradation regime is a retrospective discovery.* Its definition includes a gate that we *recompute* on the test split, so the seven cells are identified using data we have already looked at. That establishes existence; it is not a prospective validation of the criterion, and it says nothing about how often the regime occurs or how well the criterion would detect it in advance.

A pre-registered prospective test. Because (B) is the weaker claim, we measure it rather than caveat it. We selected **eight new cells**—two new model sizes, a new dataset, and horizons already in the matrix—that played no part in defining the regime, computed each cell’s gate from *validation* windows alone, and **committed the predictions to git before any condition B or D run existed**. The guarantee is structural, not procedural: the outcome needs forg_B and forg_D from runs that did not yet exist, and every quantity behind a prediction comes from train/validation only. Two rules were registered:

- **gate rule** — $R_{\text{task}}^2 \geq 0.20$ on validation $\Rightarrow$ at risk. The paper’s own criterion, pre-specified in §2.
- **dataset rule** — $\text{dataset} \in \{\text{ETTh1}, \text{Weather}\} \Rightarrow$ degradation. This is *post hoc*: among the published Moirai cells, dataset separates the regime perfectly while gate values overlap. We registered it as a *competitor*, not as a contribution, because a rule we noticed after the fact outperforming the one we published is itself evidence about the criterion. The two disagree on four of the eight cells, which is what makes the test informative rather than confirmatory.

Cell (new)	gate _{val}	forg _B	B–D held-out	Degradation?	Predicted
base Weather ($h=96$)	+0.41	−11.7 (0/3)	+9.4±2.0	no	gate ×, data ×
base Weather ($h=192$)	+0.59	−18.1 (0/3)	+12.5±8.7	no	gate ×, data ×
base ETTh2 ($h=96$)	+0.30	−13.1 (0/3)	+7.0±2.4	no	gate ×, data ✓
base ETTh2 ($h=192$)	+0.26	−8.4 (0/3)	+1.3±3.4	no	gate ×, data ✓
small Electricity7 ($h=96$)	+0.71	−16.0 (0/3)	−3.2±1.5	no	gate ×, data ✓
small Electricity7 ($h=192$)	+0.76	−18.6 (0/3)	−2.6±2.1	no	gate ×, data ✓
large ETTh1 ($h=96$)	+0.35	+14.1 (3/3)	+28.2±4.9	yes	gate ✓, data ✓

The criterion is a base-rate predictor. It flagged all seven scored cells; one degraded. That is **precision 0.14 at recall 1.00**—no false negatives, six false positives. The dataset rule scores 5/7,

496 better but not a rule either: it is correct wherever it predicted *no* degradation and wrong on both
 497 Moirai-Base/Weather cells. Taken with the published cells, degradation appears on ETTh1 at Small,
 498 Base and Large, but on Weather only at Small—so it tracks dataset crossed with model size, and
 499 **we do not offer a prospective rule**. This is the clearest support we have for §2’s “inclusion, not
 500 prescription”: the criterion says where an intervention is *interpretable*, not where fine-tuning will
 501 hurt.

502 Two further observations. The highest validation gates in the entire matrix (0.712, 0.764, Moirai-
 503 Small/Electricity7) belong to cells where adaptation *helps* or the comparison is directional—if the
 504 gate carried risk signal, those are where it should have appeared. And forg_B is negative in six of the
 505 seven new cells with *zero* positive seeds anywhere, so the regime is rarer than seven-of-30 makes it
 506 sound.

507 **A correction to the gate itself.** Auditing the above surfaced a window mismatch in the Moirai arm of
 508 `gate_all_cells.py`: the linear denominator was built with an extended lookback of 2×96 for every
 509 cell, while `finetune_forecasting.py` evaluates with lookback $+h$. The two agree at $h=96$ and
 510 diverge at $h=192$, where the numerator and denominator of R_{task}^2 were being averaged over different
 511 target windows. Corrected, six $h=192$ gate values move (largest, Moirai-Base/ETTh2 $0.72 \rightarrow 0.55$;
 512 Moirai-Small/Weather $0.80 \rightarrow 0.83$). **No cell changes gate status and the degradation set is**
 513 **unchanged**; the values in Table 1, Table 6 and App. M are the corrected ones.

514 **Protocol.** Train on train, select on validation, report on test. Nothing is selected on the test windows:
 515 condition B’s early stopping is untouched, and the ridge penalty is still the argmin of the *validation*
 516 grid, merely re-scored at that α on test. The Chronos runs share one fixed test subsample across
 517 seeds and one seeded zero-shot reference per dataset, since B–D divides by a shared denominator.
 518 Condition B was re-trained because only the encoder, not the head, had been checkpointed; its
 519 validation losses reproduce the published values to 0.008–0.039% (MPS run-to-run float noise), and
 520 all six frozen-encoder ridge optima reproduce exactly and remain interior.

521 **Which cells count as degradation.** The definition—gate-passing, $\text{forg}_B > 0$ in every seed, $\text{forg}_D < 0$
 522 in every seed—selects exactly seven of the 30 cells, and `cell_matrix.py` applies it programmati-
 523 cally so the list cannot be curated by hand. The four TimesFM cells are put through the same test and
 524 none passes it. One further cell (Moirai-Small/ETTh2, $h=192$) meets the inequalities on *means* but
 525 its harm term is only $+2.0$ with 5/10 seeds positive, so it is excluded and reported as directional. An
 526 earlier version of the paper listed four, labelling the two Moirai-Small/ETTh1 cells *preserve* by hand
 527 despite their satisfying every clause.

528 Table 6 is emitted by `scripts/cell_matrix.py`, the same function that computes the body’s
 529 numbers, rather than transcribed. An earlier version of this appendix listed four cells while Table 1
 530 already quoted held-out values for all 19—the appendix had not been updated when the re-scoring
 531 was extended, and the mismatch reads as the body overclaiming. Generating it is the only fix that
 532 stays fixed. All dispersions here and in the body are SEM; the earlier table mixed SEM with SD,
 533 which made two cells look better separated than they were.

534 **What moves.** Condition B is stable across the two segments; condition D is not. On Chronos/ETTh1
 535 the frozen head goes from 1.58/1.36/1.37 on validation to 1.81/2.51/2.07 on test, while B moves by
 536 a few percent. Chronos/ETTh2, the one cell that does not reverse, is also the one whose frozen head
 537 transfers cleanly ($1.133 \rightarrow 1.135$). So the cell where freezing genuinely cost was robust to the choice
 538 of windows, and the cells where freezing appeared to *help* were the ones that did not survive it.

539 **A consequence for the closed-form argument.** We previously justified the ridge head as a closed-
 540 form optimum that admits no epoch budget. It is an optimum *on the validation windows*. Out
 541 of sample a weak validation-selected α transfers poorly, and the protocol-matched early-stopped
 542 AdamW head—which agrees with ridge to ~ 1 pp on validation—beats it on test (Chronos/ETTh1
 543 seed 43: 2.505 against 1.957). We pre-committed to reporting the estimator validation preferred,
 544 which is ridge in every seed, with AdamW as a sensitivity; taking the better of the two on test would
 545 reintroduce exactly the fault this appendix removes. Under the AdamW leg the ETTh1 reversal is
 546 smaller but intact.

547 **III.** This cell also required its condition D to be re-run from scratch, for a different reason: the frozen
 548 control had never been trained (App. C). With a real D its validation B–D is -62.6 ± 1.4 rather than
 549 the -80.7 we reported, and -38.6 ± 1.3 on held-out windows. Its D reaches CKA 0.76–0.90, not the
 550 1.000 the earlier version asserted, because the projections feeding the encoder stay trainable.

Table 6: Every cell, both scorings. Gate and B–D are computed on the same test windows; ^f marks the three cells that gate-fail there. “Reverses” = the sign of B–D differs between the two scorings.

Cell	Seeds	Gate R^2_{task}	B–D validation	B–D held-out	Reverses?
Chronos / ETTh1 ($h=24$)	3	+0.389	+6.8±2.9	− 39.2±13.9	yes
Moirai-S / ILI ($h=24$)	10	−0.033 ^f	−62.6±1.4	− 38.6±1.3	no
Chronos / ETTm2 ($h=24$)	3	+0.651	−24.1±9.8	− 29.5±6.4	no
Moirai-B / ETTh2 ($h=96, n=1000$)	3	+0.287	+3.8±4.1	− 22.3±2.9	yes
Chronos / Weather ($h=24$)	3	+0.494	−22.4±4.5	− 21.6±3.2	no
Moirai-L / ETTh2 ($h=96, n=500$)	3	+0.129 ^f	+7.2±5.8	− 20.8±1.4	yes
Moirai-B / ETTh2 ($h=192, n=1000$)	3	+0.554	−3.9±6.2	− 20.2±1.1	no
Chronos / ETTh2 ($h=24$)	3	+0.456	−19.9±3.2	− 12.9±0.6	no
TimesFM / ETTm2 ($h=24$)	3	+0.703	−12.2±12.4	− 7.7±2.6	no
TimesFM / ETTh2 ($h=24$)	3	+0.472	+0.2±2.9	− 7.1±3.1	yes
Moirai-S / Electricity7 ($h=96$)	3	+0.535	−10.4±1.1	− 3.2±1.5	no
Chronos / Electricity ($h=24$)	3	+0.138 ^f	−10.3±3.2	− 2.7±1.9	no
Moirai-S / Electricity7 ($h=192$)	3	+0.650	−8.2±1.9	− 2.6±2.1	no
TimesFM / Weather ($h=24$)	3	+0.506	−4.8±0.6	− 1.6±1.5	no
TimesFM / ETTh1 ($h=24$)	3	+0.401	−5.4±0.9	− 1.4±2.8	no
Moirai-B / ETTm2 ($h=192$)	3	+0.714	+3.5±1.5	+ 1.3±3.4	no
Moirai-S / ETTm2 ($h=96, n=1000$)	5	+0.568	+6.5±2.9	+ 2.6±3.4	no
Moirai-S / ETTh2 ($h=96, n=500$)	10	+0.331	+10.6±2.8	+ 2.8±3.7	no
Moirai-S / ETTm2 ($h=192, n=1000$)	5	+0.722	+8.6±3.3	+ 4.7±3.2	no
Moirai-S / ETTh2 ($h=192, n=500$)	10	+0.670	+17.4±4.3	+ 6.1±3.4	no
Moirai-B / ETTm2 ($h=96$)	3	+0.588	+3.9±3.7	+ 7.0±2.4	no
Moirai-B / Weather ($h=96$)	3	+0.485	+19.3±0.8	+ 9.4±2.0	no
Moirai-B / Weather ($h=192$)	3	+0.752	+30.6±12.0	+ 12.5±8.7	no
Moirai-S / ETTh1 ($h=192, n=1000$)	5	+0.719	+14.8±1.9	+ 13.5±1.7	no
Moirai-S / ETTh1 ($h=96, n=1000$)	5	+0.578	+16.1±1.5	+ 22.1±2.3	no
Moirai-L / ETTh1 ($h=96$)	3	+0.506	+13.4±2.1	+ 28.2±4.9	no
Moirai-B / ETTh1 ($h=96$)	3	+0.505	+15.8±4.6	+ 37.5±8.2	no
Moirai-B / ETTh1 ($h=192$)	3	+0.682	+15.6±2.3	+ 47.0±13.8	no
Moirai-S / Weather ($h=96$)	3	+0.526	+55.4±4.3	+ 53.1±10.2	no
Moirai-S / Weather ($h=192$)	3	+0.825	+70.2±12.3	+ 62.1±13.6	no

551 **Scope and denominators.** Every cell in Table 6 is scored on a chronologically disjoint test split. The
552 zero-shot test denominator matters more than it looks: B–D divides by a shared zero-shot term, and
553 each run stores one measured on *validation*, so using it against a test-window difference mixes scales—
554 on Moirai-Small/ETTh2 $h=96$ the test zero-shot is 0.492 against validation’s 0.129, which inflates
555 held-out B–D by $\approx 3.8\times$. The test references come from three run sets (results/v41_zs_test,
556 v43_moirai_matrix/*/condition_A, v39_moirai_zs_test), 13 in total, and the Chronos cells
557 share one fixed test subsample and one seeded zero-shot per dataset across seeds. Two rows of Table 1
558 in an earlier version—Chronos/M4 and the TimesFM $h=96$ pair—have no paired B/D artifacts with
559 test MSEs retained and are not among the 30; they appear only in App. E. The four TimesFM $h=24$
560 cells are a different matter: they were run for this revision and scored on test windows from the outset,
561 so nothing about them was re-scored.

562 H The Strict-Freeze Control

563 Condition D freezes the encoder’s own weights but leaves the input projection (in_proj, 247,680
564 parameters) and mask_encoding (384) trainable alongside the distribution head (param_proj,
565 2,956,800). What the encoder *receives* therefore keeps changing, so its output is not a fixed function
566 of the input—which is why D’s measured CKA is 0.76–0.90 on ILI and 0.97–0.99 on Moirai-Base
567 rather than 1.000. A reader can reasonably object that “freezing wins” in the degradation cells is
568 partly the input projection re-fitting rather than the encoder being held still.

569 Condition H removes the objection: in_proj and mask_encoding are frozen too, so only
570 param_proj trains and the encoder’s output is pinned by construction (measured CKA 1.0000
571 exactly, as it must be). B–H is then a clean read of what allowing *any* encoder-side change was

worth, and B–D versus B–H separates encoder adaptation from input re-fitting. We fixed the reading rule before running: the claim is that the sign and reading of B–D survive strict freezing, so B–H is reported beside B–D on every cell, and divergence on any cell is reported as that cell’s gap being partly input re-fitting rather than being buried.

Result. All six degradation cells known at the time, plus ILI, were run; nothing was dropped. Measured CKA is 1.0000 everywhere, confirming the construction. **The reading survives on every cell:** B–H keeps B–D’s sign in all seven, and the magnitude changes by -5.8 to $+0.5$ pp—so at most 5.8 pp of a gap as large as 62 pp is input-projection re-fitting. On ILI, where adaptation helps, strict freezing makes adaptation look *more* valuable ($-38.6 \rightarrow -43.6$), which is the same direction: holding more of the model still costs more where the encoder was earning its keep. B–D is averaged over the seeds condition H ran, so Moirai-Small/ETTh1’s entries here are 3-seed values against the 5-seed values quoted in Table 1.

Cell	Seeds	B–D	B–H	for _H	CKA _H
Moirai-S / Weather ($h=192$)	3	$+62.1 \pm 13.6$	$+56.7 \pm 13.5$	-10.6 ± 0.3	1.0000
Moirai-S / Weather ($h=96$)	3	$+53.1 \pm 10.2$	$+50.1 \pm 6.6$	-12.8 ± 0.6	1.0000
Moirai-B / ETTh1 ($h=192$)	3	$+47.0 \pm 13.8$	$+47.5 \pm 13.7$	-8.0 ± 0.4	1.0000
Moirai-B / ETTh1 ($h=96$)	3	$+37.5 \pm 8.2$	$+32.5 \pm 8.2$	-9.7 ± 0.4	1.0000
Moirai-S / ETTh1 ($h=96, n=1000$)	3	$+23.2 \pm 3.6$	$+20.4 \pm 3.6$	-0.6 ± 0.2	1.0000
Moirai-S / ETTh1 ($h=192, n=1000$)	3	$+14.5 \pm 1.4$	$+8.7 \pm 1.5$	$+1.2 \pm 0.4$	1.0000
Moirai-S / ILI ($h=24$)	10	-38.6 ± 1.3	-43.6 ± 0.7	-7.1 ± 0.5	1.0000

One clause does not survive, and it is not the headline one. “Degradation” has two parts: fine-tuning ends up worse than the un-tuned model, *and* freezing improves on it. The first part is a property of condition B and is untouched by this control. The second is condition D’s, and on Moirai-Small/ETTh1 $h=192$ the strictly-frozen model does *not* beat the un-tuned model: $\text{for}_H = +1.2 \pm 0.4\%$, positive in all 3 seeds, where for_D is $-4.6 \pm 0.6\%$. On that cell the head alone cannot recover the zero-shot forecast; D’s improvement came partly from re-fitting the input projection. The other five cells keep both parts (for_H from -0.6% to -12.8%), and B–H stays positive on all six, so “freezing beats full fine-tuning” holds throughout while “freezing also beats not fine-tuning” holds strictly on five of six. We report the cell rather than re-defining the term around it.

I CKA Pooling Robustness, and the Re-Analysis that Retracts the Probe Axis

A consequence worth stating plainly: because the magnitude of CKA depends on how the representation is constructed, this paper’s claims are strongest where they are *ordinal and interventional*—which cells drift, and what freezing costs on each—and weakest if read as statements about absolute CKA values. The drift *flag* is stable under pooling; the number attached to it is not.

Pooling robustness. CKA is reported on mean-pooled encoder outputs throughout. Because mean-pooling discards token-level structure, we recomputed whole-encoder CKA on the same ten fine-tuned ETTh2 encoders ($n=10k$, $h=96$, condition B) both mean-pooled and token-flattened, over 500 held-out 192-timestep windows.

Representation	CKA (mean $\pm$ SD, 10 seeds)	Seed range
Mean-pooled (as reported)	0.561 ± 0.157	0.328–0.840
Token-flattened	0.732 ± 0.036	0.685–0.797

Two things follow. The *qualitative* drift flag is robust: both values sit far below the 0.95 threshold, so every cell keeps its assignment. The *magnitude* is not: token-level similarity is substantially higher and far less seed-variable (± 0.036 vs ± 0.157), so part of the drift reported as ~ 0.52 is an artifact of pooling. We regard this as consistent with, rather than damaging to, the thesis—it is a further reason not to read a single scalar as a measure of damage—but CKA magnitudes should be read as pooling-dependent.

Probe protocol, recovered and verified. The probe pipeline was not documented in enough detail to re-run, so we reconstructed it from source and verified it against the published numbers. The protocol is: the encoder input window is **192 steps** (extended lookback = lookback + horizon), the probe

613 **trains on the first 300 validation windows and is evaluated on the first 200 test windows**, the
614 target is the next 96 steps across **all seven features** flattened, and the probe is Ridge($\alpha=1.0$) on
615 mean-pooled encoder outputs. Under exactly this protocol the trained probe reproduces:

	Trained 96-step probe	$R^2(\text{PT})$	$R^2(\text{FT})$	ΔR^2 (10 seeds)
616	Published	-6.90	-6.23	+0.67 $\pm$ 0.23, 0/10 neg
	This reproduction	-6.9043	-6.2364	+0.6679 $\pm$ 0.2262, 0/10 neg

617 The headline ΔR^2 therefore reproduces to four significant figures from the released checkpoints.
618 Three protocol details each move the result by an order of magnitude and are easy to get wrong: using
619 the 192-step window rather than a 96-step context, evaluating across the validation $\rightarrow$ test boundary
620 rather than within one split, and using all seven features as the target rather than the target channel
621 alone. Substituting a within-validation split, for instance, moves $R^2(\text{PT})$ from -6.90 to -0.13.

622 **A transparent re-analysis, and what it costs us.** Rather than keep searching for the target normal-
623 ization behind the published auxiliary floors, we specified a protocol with every knob fixed in advance
624 and report what it gives (`scripts/probe_transparent.py`). Windows: the same 500. Features:
625 mean-pooled encoder output, standardized on probe-train. Probe: Ridge with α chosen by 5-fold
626 CV on the *pretrained* encoder’s probe-train over $\log_{10} \alpha \in [-3, 6]$, then that same α reused for all
627 ten fine-tuned encoders so the two are never tuned apart. Statistic: ΔR^2 per seed with a percentile
628 bootstrap 95% CI over the ten seeds (10,000 resamples). Two additions decide the outcome:

- 629 • **Two splits.** T is the published temporal split (train on validation, test on test). R randomizes the
630 same 500 windows into 300/200. The pair separates decodability from the validation $\rightarrow$ test level
631 shift.
- 632 • **A permuted-target noise floor.** The identical pipeline run on 20 independently shuffled copies of
633 each target. Any real ΔR^2 must exceed the spread the pipeline produces when the target carries
634 no information at all.

Table 7: Transparent protocol, Moirai-Small/ETTh2, 10 released checkpoints. Under the published split T no probe generalizes ($R^2(\text{PT}) < 0$) and every $|\Delta R^2|$ is far below the permuted-target floor. Under R , where every probe generalizes, no $|\Delta R^2|$ exceeds 0.011 and the trained probe *decreases*.

Split	Target	$R^2(\text{PT})$	ΔR^2	95% CI	floor
T	Forecast-96 (trained)	-5.23	-0.0005	[-0.001, +0.000]	0.0005
	Mean	-164.09	+3.22	[-3.60, +10.00]	102.1
	Variance	-20.50	-8.22	[-11.94, -4.14]	132.3
	Lag-1 autocorr.	-2.71	-0.80	[-2.02, +0.16]	30.7
R	Forecast-96 (trained)	+0.678	-0.0042	[-0.0061, -0.0025]	0.019
	Mean	+0.994	-0.0014	[-0.0022, -0.0006]	0.226
	Variance	+0.937	+0.0025	[-0.0014, +0.0068]	0.169
	Lag-1 autocorr.	+0.750	+0.0103	[+0.0012, +0.0194]	0.223

635 **The published ΔR^2 is a regularization artifact.** Holding the split at T and sweeping α , the
636 trained probe’s ΔR^2 decays monotonically—+0.668, +0.271, +0.066, +0.034, +0.005, +0.000
637 at $\alpha = 1, 10, 10^2, 10^3, 10^4, 10^5$ —while $R^2(\text{PT})$ stays between -6.90 and -5.23 throughout. The
638 effect is largest exactly where the probe extrapolates worst, and CV on the training split selects
639 $\alpha=10^6$, at which it vanishes. Under split R , where the trained probe attains $R^2(\text{PT}) = +0.68$, ΔR^2
640 is at most +0.017 at any α and is negative at the CV-selected value. There is no configuration in
641 which the trained probe both generalizes and increases by anything near +0.67. (Standardizing the
642 targets is provably *not* a knob: with an intercept, Ridge is equivariant under affine target scaling, and
643 the sweep confirms the rows are identical to the last digit.)

644 **This contradicts our own stated criterion.** Section B rejects the Ridge version of the delta1 probe
645 precisely because “the zero-shot Ridge baseline is deeply negative ($R^2 \approx -2.3$), so any encoder shift
646 inflates ΔR^2 .” The headline trained probe sits at $R^2(\text{PT}) = -6.90$ —further into that regime than the
647 probe we disqualified—and we did not apply the criterion to it. Applied consistently, it disqualifies
648 the published trained/auxiliary comparison as well.

649 **What survives.** Under the one split where all seven probes generalize, every $|\Delta R^2|$ is ≤ 0.011 :
650 the trained probe decreases (-0.0042 , 10/10 seeds), mean and trend decrease, lag-1 autocorrelation
651 *increases*, and variance, lag-24 and spectral centroid are indistinguishable from zero. That is not
652 the claimed asymmetry; it is closer to no detectable change in any direction. We therefore **retract**
653 **the probe axis** as unsupported by a protocol we are willing to defend, and we have left the original
654 numbers in Table 2 rather than deleting them so the two can be compared. The frozen-encoder
655 intervention is unaffected: it shares no machinery with the probes, manipulates the encoder rather
656 than fitting a readout on it, and is measured in forecast MSE rather than R^2 . It is where the paper’s
657 weight now rests, and the two inversions it supports (§3) were never probe-dependent.

658 J Predicting the Reading Without the Frozen Control

659 An earlier version listed this as open; within Moirai it is not. Condition B’s *own* held-out forgetting
660 agrees in sign with B–D in 15 of the 21 Moirai cells, and the matched within-checkpoint contrast is
661 sharper: Moirai-Base forgets under B on ETTh1 at both horizons ($+22.8\%$, $+39.5\%$) and favours
662 freezing, while the *same* checkpoint on ETTh2 improves under B (-31.4%) and strongly favours
663 adaptation. One fine-tuning run therefore distinguishes regimes CKA does not ($\rho = +0.897$ vs CKA’s
664 $+0.181$, $n=21$).

665 We state the coupling rather than bank it. B–D $\equiv \text{forg}_B - \text{forg}_D$ identically, and forg_D barely moves
666 (SD 8.0 against forg_B ’s 26.8), so the rule works because the frozen arm is nearly constant—an
667 empirical fact about Moirai, not a theorem. It fails on Chronos ($\rho = -0.100$, limitation iii), and on
668 TimesFM, the other arm where forg_B is a clean measurement, the sign agrees 4 of 4 while forg_D
669 varies nearly as much as forg_B (SD 19.2 vs 22.0): the mechanism does not transfer even where the
670 rule does. It is reported here, not in the body, because it is a fourth result competing with the three
671 the paper is actually about, and because a rule that holds on one backbone for a reason that does not
672 generalise is not a contribution we want weighed alongside them.

673 K Related Work

674 Catastrophic-forgetting mitigations—EWC [Kirkpatrick et al., 2017], LwF [Li and Hoiem, 2018],
675 LoRA [Hu et al., 2022], L2-SP [Li et al., 2018]—are designed to limit how far a fine-tuned model
676 moves from its pretrained state, but they are typically applied without first asking whether the
677 pretrained state is worth preserving on the target task. Our pretrained-capability inclusion criterion
678 is a downstream-only check for exactly that question, and the preservation spectrum (Appendix D)
679 shows why it matters: methods that raise representational similarity do not reliably improve utility,
680 and the ordering reverses with horizon. The reflex itself is well founded where it originated: Luo
681 et al. [2023] document substantial forgetting in LLMs under continual fine-tuning. Our claim is not
682 that the phenomenon is absent but that on TSFMs the drift measurement used to detect it does not
683 order downstream harm.

684 Closest to this work, Wiliński et al. [2025] study TSFM representations directly: they measure layer
685 self-similarity within and across model sizes, find block-like redundancy that supports informed
686 pruning, and steer the latent space to inject concepts such as periodicity and trend. Their interventions
687 act on a *trained model’s activations* to characterize and control what the representation encodes;
688 ours acts on the *learning process*—freezing the encoder and re-running fine-tuning—to ask whether
689 representation change was necessary for the downstream gain. The two are complementary, and their
690 redundancy finding is a plausible mechanism for what we observe: if TSFM representations carry
691 substantial redundancy, a large measured drift can be reorganization within that slack rather than
692 loss of task-relevant structure, which is exactly the gain-independent restructuring our intervention
693 isolates. Neither their analysis nor ours predicts a cell’s B–D without running both conditions.

694 Kumar et al. [2022] show that fine-tuning can distort pretrained features and that linear probing
695 followed by fine-tuning mitigates it; Neyshabur et al. [2020] study what transfers between pretrained
696 and fine-tuned networks. Neither provides a pre-fine-tuning screen for when preservation is warranted,
697 which is the gap this paper addresses. On the measurement side we use linear CKA [Kornblith et al.,
698 2019]; our pooling analysis (Appendix I) is a reminder that its magnitude is representation-dependent
699 and should be read as an operational drift detector rather than a measure of damage.

700 Pretraining overlap for the backbones we test is reported in Appendix E.

L Moirai-Base Frozen-Encoder Run

The Moirai-Base/ETTh2 cell is the most heavily drifted cell in the study but originally had no frozen-encoder control, so nothing was known about what encoder adaptation contributed there. We ran condition D for this revision under the published condition-B protocol (Moirai-Base, ETTh2, $h=96$, $n=10k$, lr 10^{-4} , early stopping), on the same ten seeds as the published condition B, on MPS.

Seed	Forgetting%	Final CKA	Weight drift	Best epoch
42	-49.33	0.984	1.88	7
101	-54.13	0.972	2.87	16
123	-45.85	0.985	1.90	6
202	-40.14	0.983	1.77	6
303	-43.27	0.990	1.44	4
456	-40.71	0.979	2.58	12
777	-38.03	0.977	2.15	9
789	-49.86	0.983	2.13	8
888	-41.25	0.969	2.53	10
999	-50.71	0.976	2.50	11
Mean $\pm$ std	-45.33 $\pm$ 5.43	0.980 $\pm$ 0.006		

All 10 seeds are negative (SEM 1.72). Against the published condition B ($-54.2\pm 8.6\%$, 10 seeds) this is a B-D gap of 8.9 pp: on this cell encoder adaptation supplies a modest fraction of the gain, despite CKA falling to 0.25–0.33 under full fine-tuning—the deepest drift anywhere in the study. The first three seeds we ran (42, 101, 123) averaged -49.8% and gave a 4.4 pp gap, reported in an earlier version; extending to the full ten seeds moved D up to -45.3% and roughly doubled the gap, so that earlier estimate rested on a favourable subset. The remaining caveat is hardware: the published B ran on CUDA and this D on MPS. At ten seeds each the standard deviations still overlap, but the means no longer do (Welch $t=2.8$ on the 8.9 pp difference), so this cell is better read as a small-but-real adaptation contribution than as $B\approx D$.

A matched, paired comparison at $n=1000$. A separate Moirai-Base B/D pair exists at $n=1000$ (20 epochs, no early stopping) with both conditions run on the *same* seeds and hardware, which the $n=10k$ comparison above cannot offer (its B is the published CUDA run, its D the 10-seed MPS run reported here).

Seed	B (full FT)	D (frozen)	B-D
42	-35.33	-41.64	+6.31
123	-47.93	-43.62	-4.31
456	-47.57	-56.88	+9.31
Mean	-43.61	-47.38	+3.77

Freezing the encoder is if anything *better* here, in 2 of 3 seeds. Because the seeds and stack match, this is not exposed to the cross-hardware caveat above, and it reaches the same conclusion at a different n : on Moirai-Base, encoder adaptation does not account for the gain.

The run also shows that condition D is not representation-frozen. Early stopping selected epochs 6–16 rather than epoch 0, and the frozen-encoder CKA settles at 0.97–0.98 with non-zero weight drift (1.88–2.87), because the input/output projections outside the encoder remain trainable. In cells where early stopping selects the untrained state (ILI, Chronos/M4) the measured CKA is 1.000, which is what earlier drafts described—incorrectly—as holding by construction.

M Threshold Sensitivity

Both thresholds are operational decision boundaries for this study, not universal constants, and neither was tuned to produce the reported signatures.

Preservation-priority gate (20%), and a margin claim we withdraw. Earlier versions reported the observed gate values as bimodal with a wide margin—“no cell falls in $[0, 28)\%$, so any threshold in

10–30% produces identical assignments”. That was read on the windows model selection uses. Re-computed on the test windows for all 30 cells from one code path (`scripts/gate_all_cells.py`, §3), the values are −3.3, 12.9, 13.8, 28.7, 33.1, 38.9, 40.1, 45.6, 47.2, 49.4, 50.5, 50.6, 52.6, 55.4, 56.8, 57.8, 65.1, 67.0, 68.2, 70.3, 71.9, 72.2, 82.5% (the four TimesFM cells contribute 40.1, 47.2, 50.6 and 70.3). Three cells fall below the threshold (−3.3, 12.9, 13.8) and one sits just above it at 28.7%, so the distribution is not bimodal out of sample and the threshold is *not* margin-free: moving it to 30% would additionally reclassify Moirai-Base/ETTh2 $h=96$, and moving it to 10% would restore Chronos/Electricity and Moirai-Large/ETTh2 to gate-pass. The six degradation cells are unaffected across the whole 10–50% range (they sit at 50.5–82.5%), which is the assignment the paper’s claims depend on; the reclassified cells all read *adapt*, where the gate is not what licenses the reading.

CKA threshold reclassification. The gate analysis above shows wide margins; the CKA cutoff deserves the same treatment. We reclassify every cell at four thresholds using the observed CKA values (ranges span seeds, not datasets).

Thresh.	Moirai-B (0.25–0.33)	Moirai/ILI (0.456)	Chronos/ETT (0.087–0.179)	TimesFM ($h=24$) (0.19–0.64)	Chronos/M4 (0.98–0.999)	TimesFM ($h=96$) (0.995–0.999)
0.900	drift	drift	drift	drift	stable	stable
0.925	drift	drift	drift	drift	stable	stable
0.950	drift	drift	drift	drift	stable	stable
0.975	drift	drift	drift	drift	stable	stable

Every assignment is identical at all four thresholds, and no cell splits. Note that TimesFM appears twice and on opposite sides: its $h=96$ cells are the most stable in the study and its $h=24$ cells among the most drifted, so “stable” is not a property of the backbone. The margin is not symmetric, and it is worth being explicit about which end is tight: the drifting cells sit at least 0.26 below the lowest cutoff considered, whereas the stable cells clear the highest one by only 0.005 (Chronos/M4, whose lowest seed is 0.98) and 0.020 (TimesFM at $h=96$). So it is the stable end, not the drifting end, that would become threshold-sensitive if the cutoff were pushed past 0.975. We state the operating range explicitly rather than claiming threshold-independence: the conclusions hold for cutoffs in $[0.90, 0.975]$.

Earlier versions of this table listed the Chronos/ETT cells at 0.886–0.949 and therefore reported them as *split* below 0.95. That range was misattributed—it belongs to the Moirai-Small sample-size sweep, not to Chronos—and the measured Chronos values are an order of magnitude lower, which removes the split.

CKA (0.95). Observed CKA clusters away from the cutoff at the drifting end: drifting cells sit at 0.087–0.179, 0.19–0.64, 0.25–0.33, and 0.456; stable cells at 0.98–0.999. A random-initialization baseline gives CKA 0.000 ± 0.000 (5 seeds), fixing the floor of the scale, so the Chronos/ETT cells at 0.087–0.179 are much nearer that floor than the threshold. This is also why the Chronos/ETTh1-vs-ETTh2 comparison cannot be resolved by CKA at any cutoff: both cells are unambiguously drifted on every reading, yet the less-drifted of the two is the one that loses more to freezing.